

# MapReduce

**Thoai Nam**

High Performance Computing Lab (HPC Lab)  
Faculty of Computer Science and Engineering  
HCMC University of Technology

# Ref

- MapReduce algorithm design, Jimmy Lin



processes 20 PB a day (2008)  
crawls 20B web pages a day (2012)



>10 PB data, 75B DB  
calls per day (6/2012)

>100 PB of user data +  
500 TB/day (8/2012)



S3: 449B objects, peak 290k  
request/second (7/2011)  
IT objects (6/2012)



640K ought to be  
enough for anybody.

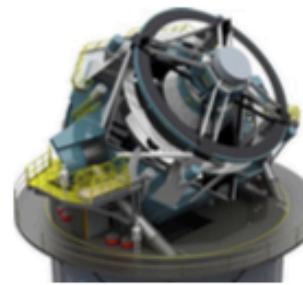


150 PB on 50k+ servers  
running 15k apps (6/2011)



Wayback Machine: 240B web  
pages archived, 5 PB (1/2013)

LHC: ~15 PB a year



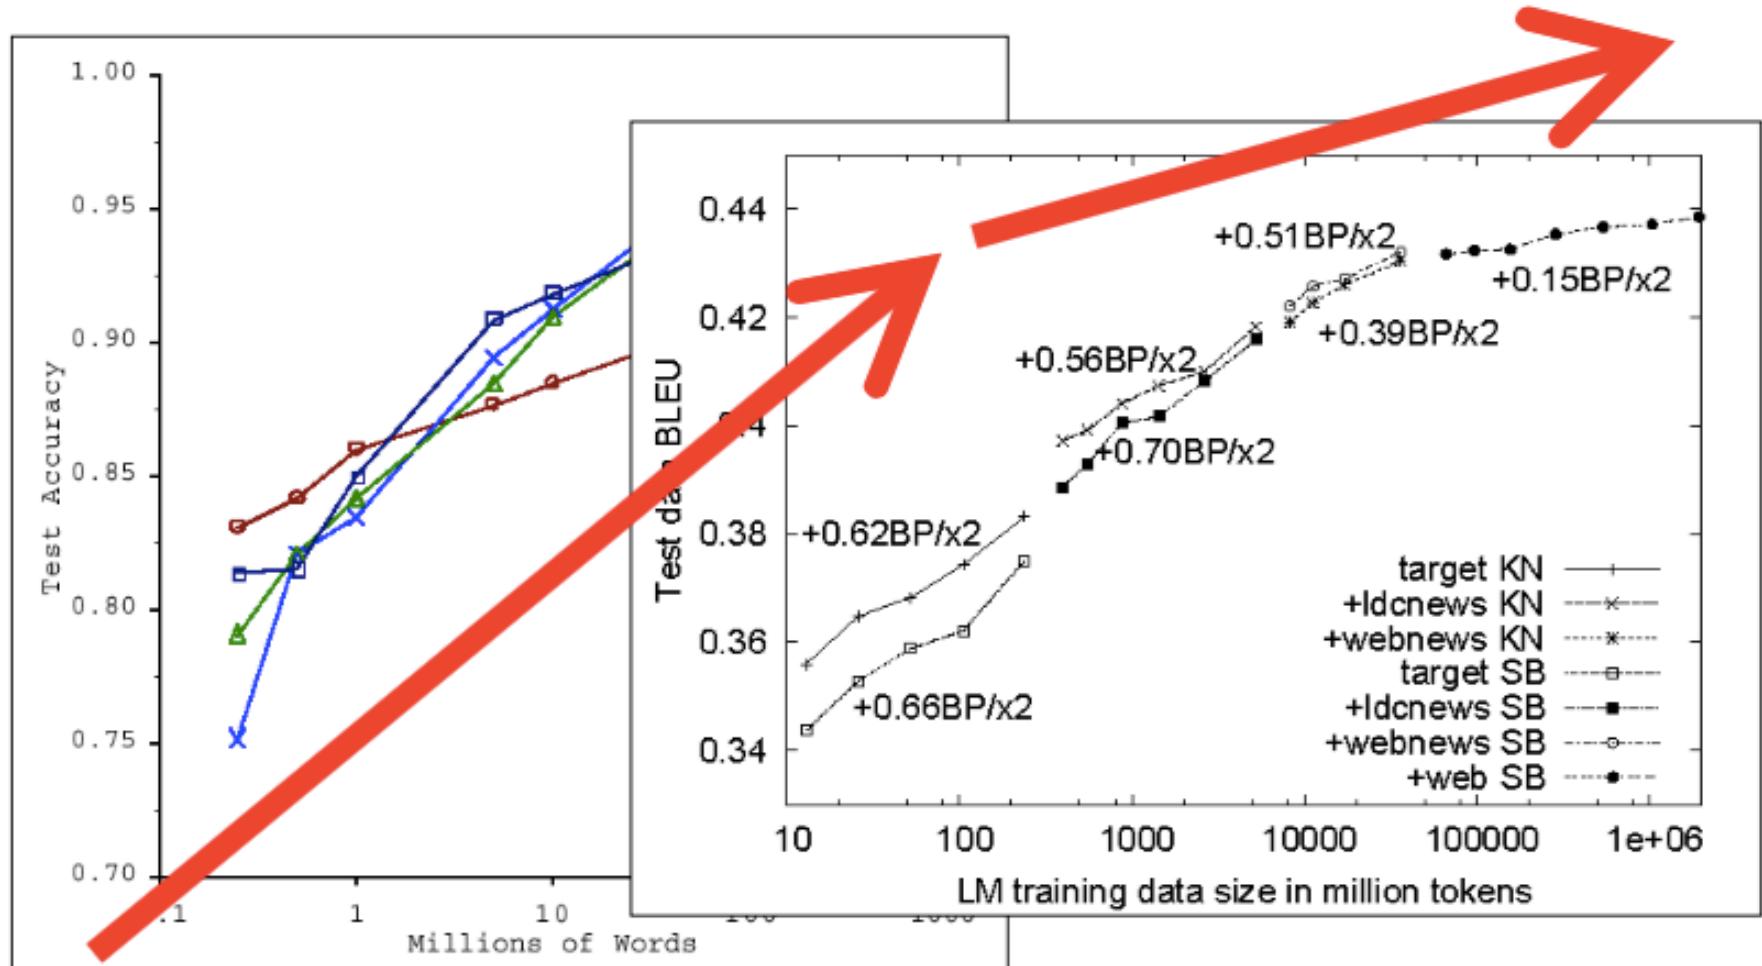
LSST: 6-10 PB a year  
(~2015)



**How much data?**

# No data like more data!

s/knowledge/data/g;





# MapReduce

# Typical Big Data Problem

- Iterate over a large number of records

**Map** Extract something of interest from each

- Shuffle and sort intermediate results
- Aggregate intermediate results
- Generate final output

**Reduce**

**Key idea: provide a functional abstraction for these two operations**

# MapReduce: A Real World Analogy

## Coins Deposit



# MapReduce: A Real World Analogy

## Coins Deposit



## Coins Counting Machine

# MapReduce: A Real World Analogy

## Coins Deposit

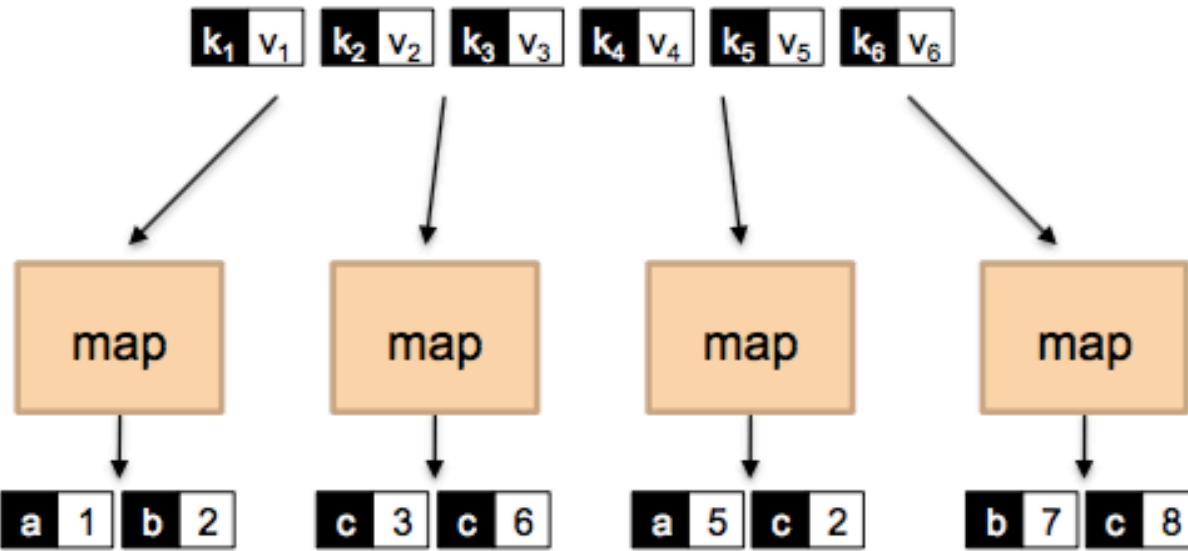


**Mapper:** Categorize coins by their face values

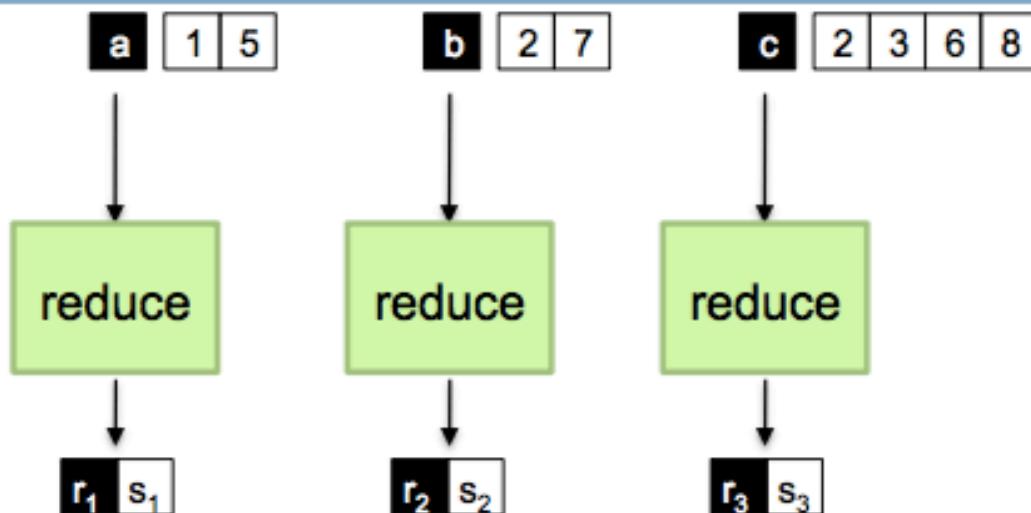
**Reducer:** Count the coins in each face value *in parallel*

# MapReduce

- Programmers specify two functions:  
**map** ( $k_1, v_1$ )  $\rightarrow$  [ $k_2, v_2$ ]  
**reduce** ( $k_2, [v_2]$ )  $\rightarrow$  [ $k_3, v_3$ ]  
(All values with the **same key** are sent to the same reducer)
- The execution framework handles everything else...



**Shuffle and Sort: aggregate values by keys**



# MapReduce

- Programmers specify two functions:

**Map** (k<sub>1</sub>, v<sub>1</sub>) → <k<sub>2</sub>, v<sub>2</sub>>\*

**Reduce** (k<sub>2</sub>, list (v<sub>2</sub>)) → list (v<sub>3</sub>)

(All values with the **same key** are sent to the same reducer)

- The execution framework handles everything else...

# MapReduce “runtime”

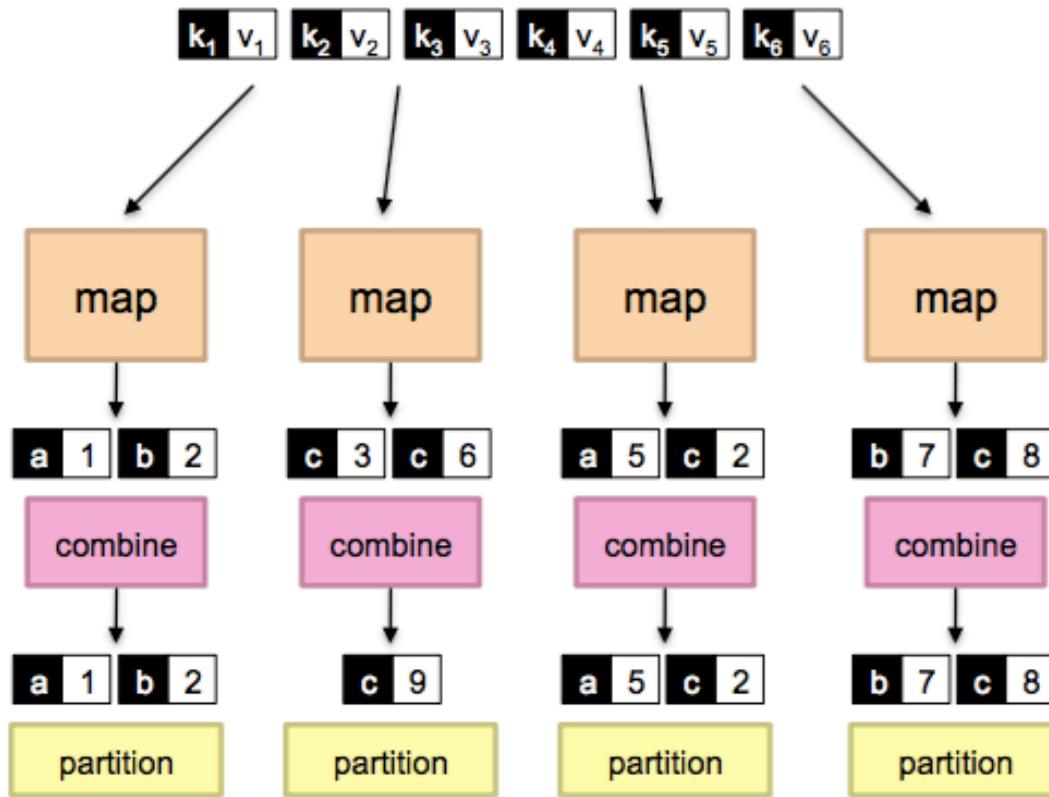
- Handles scheduling
  - Assigns workers to map and reduce tasks
- Handles “data distribution”
  - Moves processes to data
- Handles synchronization
  - Gathers, sorts, and shuffles intermediate data
- Handles errors and faults
  - Detects worker failures and restarts
- Everything happens on top of a distributed file system

# Synchronization & ordering

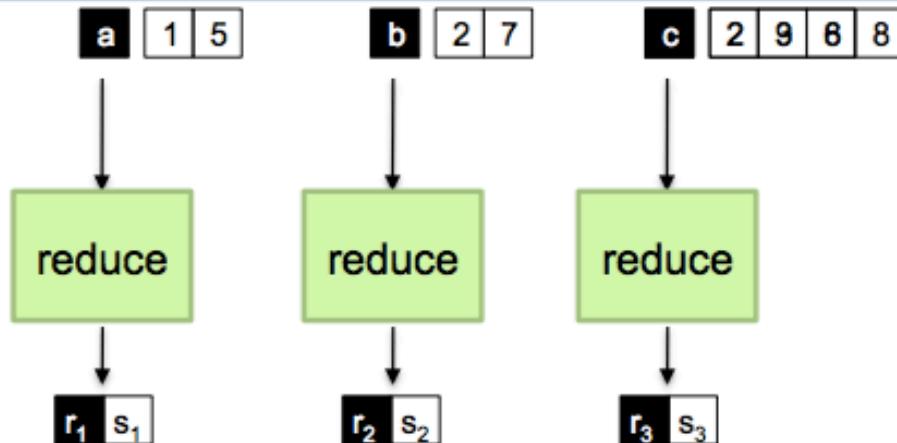
- Barrier between map and reduce phases
  - But intermediate data can be copied over as soon as mappers finish
- Keys arrive at each reducer in sorted order
  - No enforced ordering *across* reducers

# MapReduce

- Programmers specify two functions:  
**Map** ( $k_1, v_1$ )  $\rightarrow$   $\langle k_2, v_2 \rangle^*$   
**Reduce** ( $k_2, \text{list}(v_2)$ )  $\rightarrow$   $\text{list}(v_3)$   
(All values with the **same key** are sent to the same reducer)
- The execution framework handles everything else...
- Not quite...usually, programmers also specify:  
**partition** ( $k_2, \text{number of partitions}$ )  $\rightarrow$  partition for  $k_2$ 
  - Often a simple hash of the key, e.g.,  $\text{hash}(k') \bmod n$
  - Divides up key space for parallel reduce operations  
**combine** ( $k_2, v_2$ )  $\rightarrow$   $\langle k_2, v_2 \rangle^*$ 
  - Mini-reducers that run in memory after the map phase
  - Used as an optimization to reduce network traffic

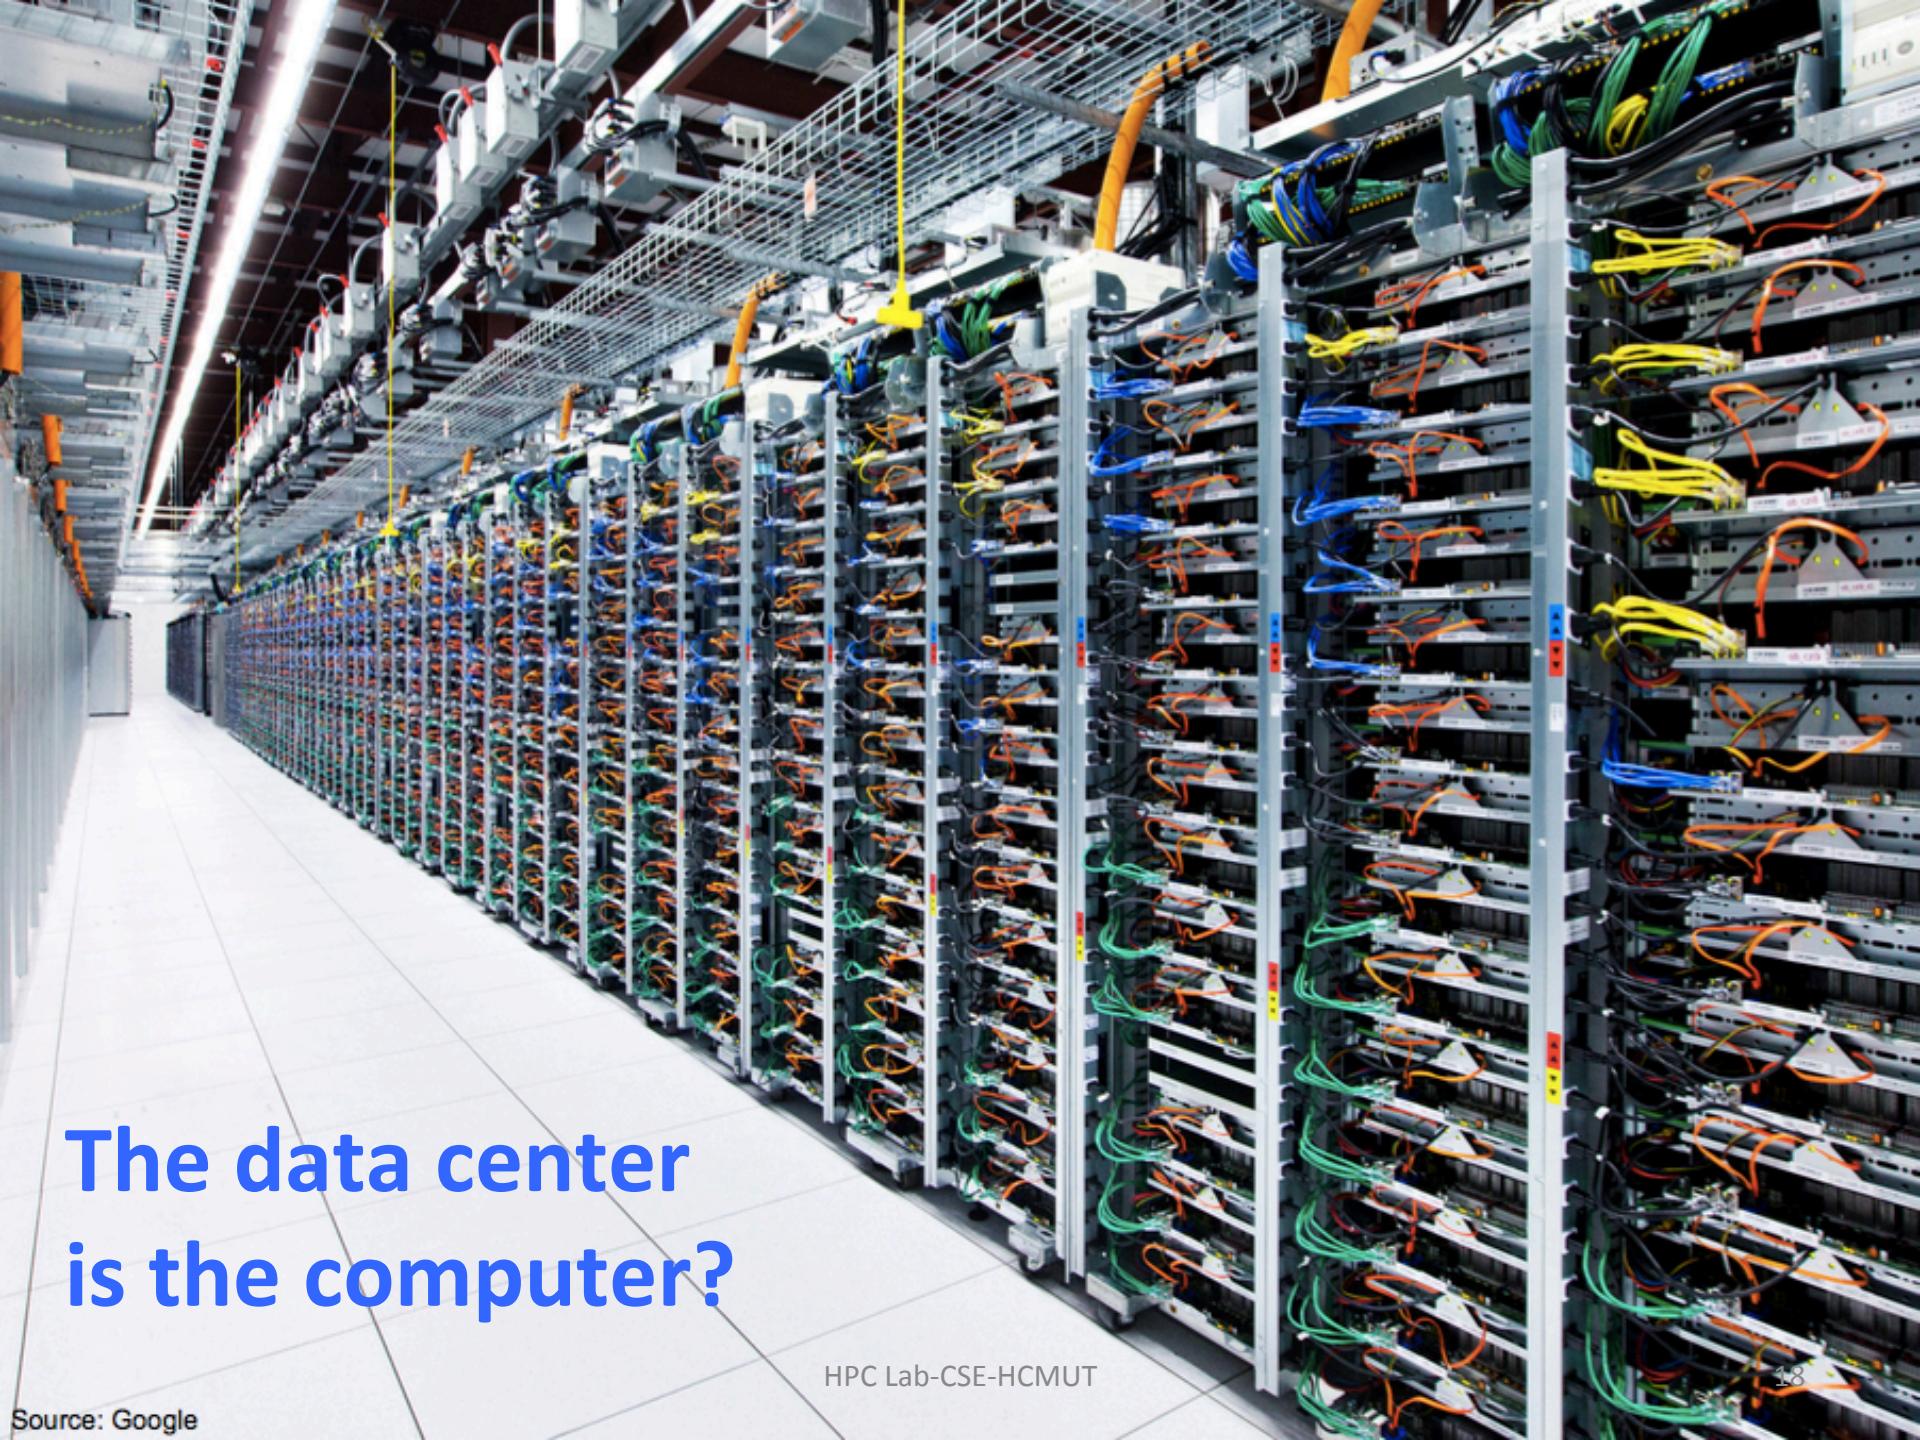


### Shuffle and Sort: aggregate values by keys



# What's the big deal?

- Developers need the right level of abstraction
  - Moving beyond the von Neumann architecture
  - We need better programming models
- Abstractions hide low-level details from the developers
  - No more race conditions, lock contention, etc.
- MapReduce separating the *what* from *how*
  - Developer specifies the computation that need to be performed
  - Execution framework (“runtime”) handles actual execution



The data center  
is the computer?

# MapReduce can refers to...

- The programming model
- The execution framework (aka “runtime”)
- The specific implementation

# MapReduce Implementations

- Google has a proprietary implementation in C++
  - Bindings in Java, Python
- Hadoop is an open-source implementation in Java
  - Development led by Yahoo, now an Apache project
  - Used in production at Yahoo, Facebook, Twitter, LinkedIn, Netflix, etc.
  - The *de facto* big data processing platform
  - Rapidly expanding software ecosystem
- Lots of custom research implementations
  - For GPUs, cell processors, etc.



# MapReduce algorithm design

- The execution framework handles “everything else”...
  - Scheduling: assigns workers to map and reduce tasks
  - “Data distribution”: moves processes to data
  - Synchronization: gathers, sorts, and shuffles intermediate data
  - Errors and faults: detects worker failures and restarts
- Limited control over data and execution flow
  - All algorithms must expressed in m,r,c,p
- You don’t know:
  - Where mappers and reducers run
  - When a mapper or reducer begins or finishes
  - Which input a particular mapper is processing
  - Which intermediate key a particular reducer is processing

# Apache Hadoop

# Data volumes: Google Example

- Analyze **10 billion web pages**
- Average size of a webpage: **20KB**
- Size of the collection:  $10 \text{ billion} \times 20\text{KBs} = 200\text{TB}$
- HDD hard disk read bandwidth: **150MB/sec**
- Time needed to **read all web pages (without analyzing them)**: 2 million seconds = more than **15 days**
- A single node architecture is not adequate

# Data volumes: Google Example with SSD

- Analyze 10 billion web pages
- Average size of a webpage: 20KB
- Size of the collection: 10 billion x 20KBs = 200TB
- SSD hard disk read bandwidth: 550MB/sec
- Time needed to read all web pages (without analyzing them): 2 million seconds = more than 4 days
- A single node architecture is not adequate

# Apache Hadoop

- Scalable fault-tolerant distributed system for Big Data
  - Distributed Data Storage
  - Distributed Data Processing
  - Borrowed concepts/ideas from the systems designed at Google (Google File System for Google's MapReduce)
  - Open source project under the Apache license
    - But there are also many commercial implementations (e.g., Cloudera, Hortonworks, MapR)

# Hadoop History

- Dec 2004 - Google published a paper about GFS
- July 2005 - Nutch uses MapReduce
- Feb 2006 - Hadoop becomes a Lucene subproject
- Apr 2007 - Yahoo! runs it on a 1000-node cluster
- Jan 2008 - Hadoop becomes an Apache Top Level Project
- Jul 2008 - Hadoop is tested on a 4000 node cluster
- Feb 2009 - The Yahoo! Search Webmap is a Hadoop application that runs on more than 10,000 core Linux cluster
- June 2009 - Yahoo! made available the source code of its production version of Hadoop
- In 2010 Facebook claimed that they have the largest Hadoop cluster in the world with 21 PB of storage
  - On July 27, 2011 they announced the data has grown to 30PB.

# Hadoop vs. HPC

- **Hadoop**
  - Designed for **Data intensive workloads**
  - Usually, no CPU demanding/intensive tasks
- **HPC** (High-performance computing)
  - A supercomputer with a high-level computational capacity
    - Performance of a supercomputer is measured in floating-point operations per second (FLOPS)
  - Designed for **CPU intensive tasks**
  - Usually it is used to process "small" data sets

# Hadoop: main components

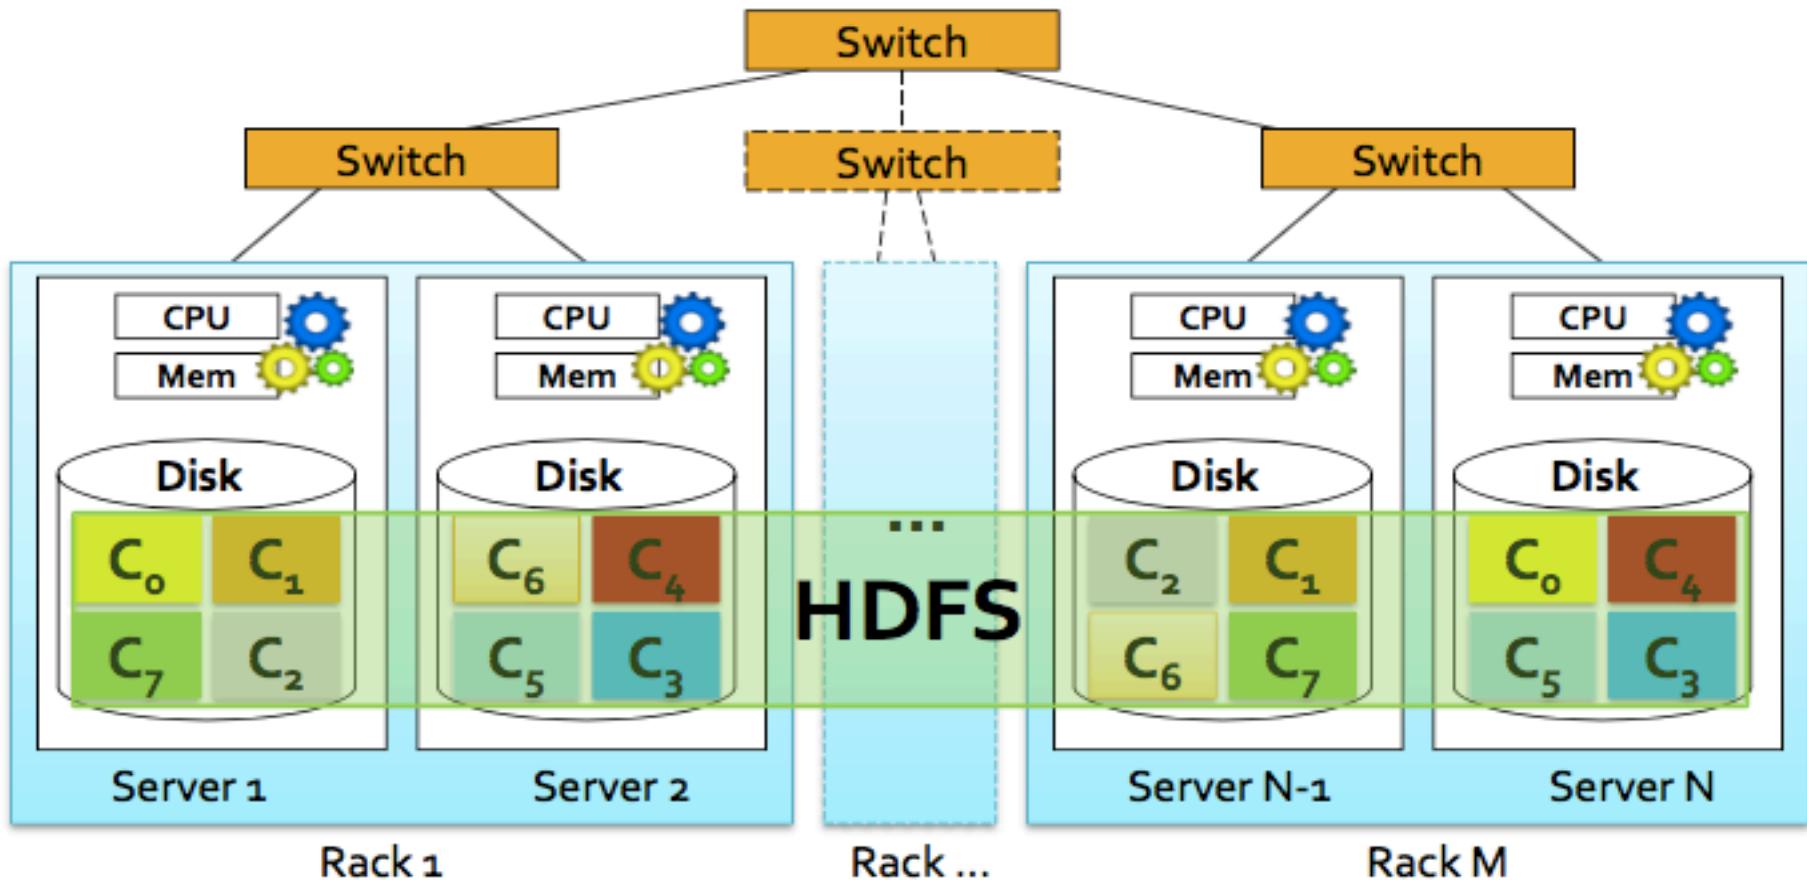
- Core components of Hadoop:
  - Distributed Big Data Processing Infrastructure based on the MapReduce programming paradigm
    - Provides a high-level abstraction view
      - Programmers do not need to care about task scheduling and synchronization
    - Fault-tolerant
      - Node and task failures are automatically managed by the Hadoop system
  - HDFS (Hadoop Distributed File System)
    - High availability distributed storage
    - Fault-tolerant

# HDFS (Hadoop File System)

# HDFS

- HDFS is a distributed file system that is fault tolerant, scalable and extremely easy to expand
- HDFS is the primary distributed storage for Hadoop applications
- HDFS provides interfaces for applications to move themselves closer to data
- HDFS is designed to ‘just work’, however a working knowledge helps in diagnostics and improvements

# HDFS: a distributed file system

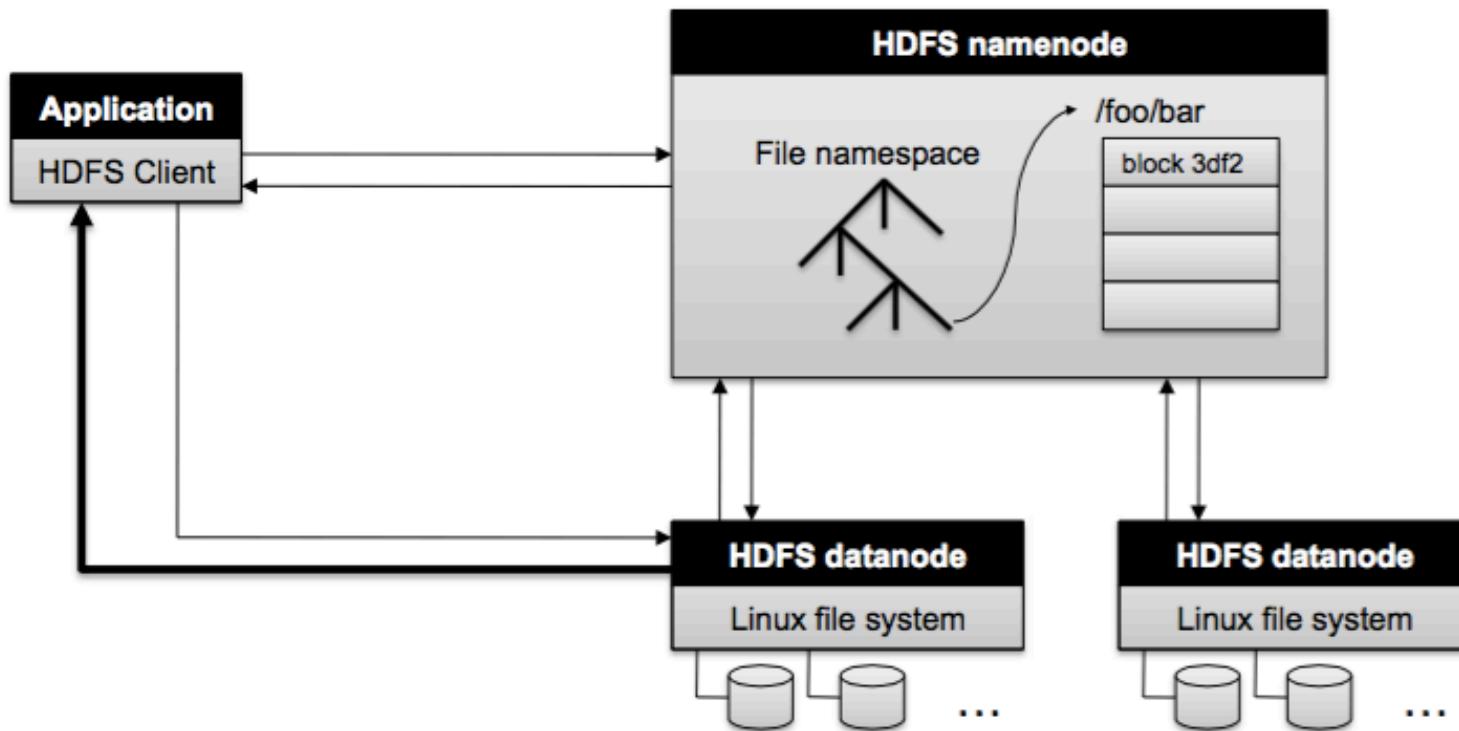


Example with number of replicas per chunk = 2

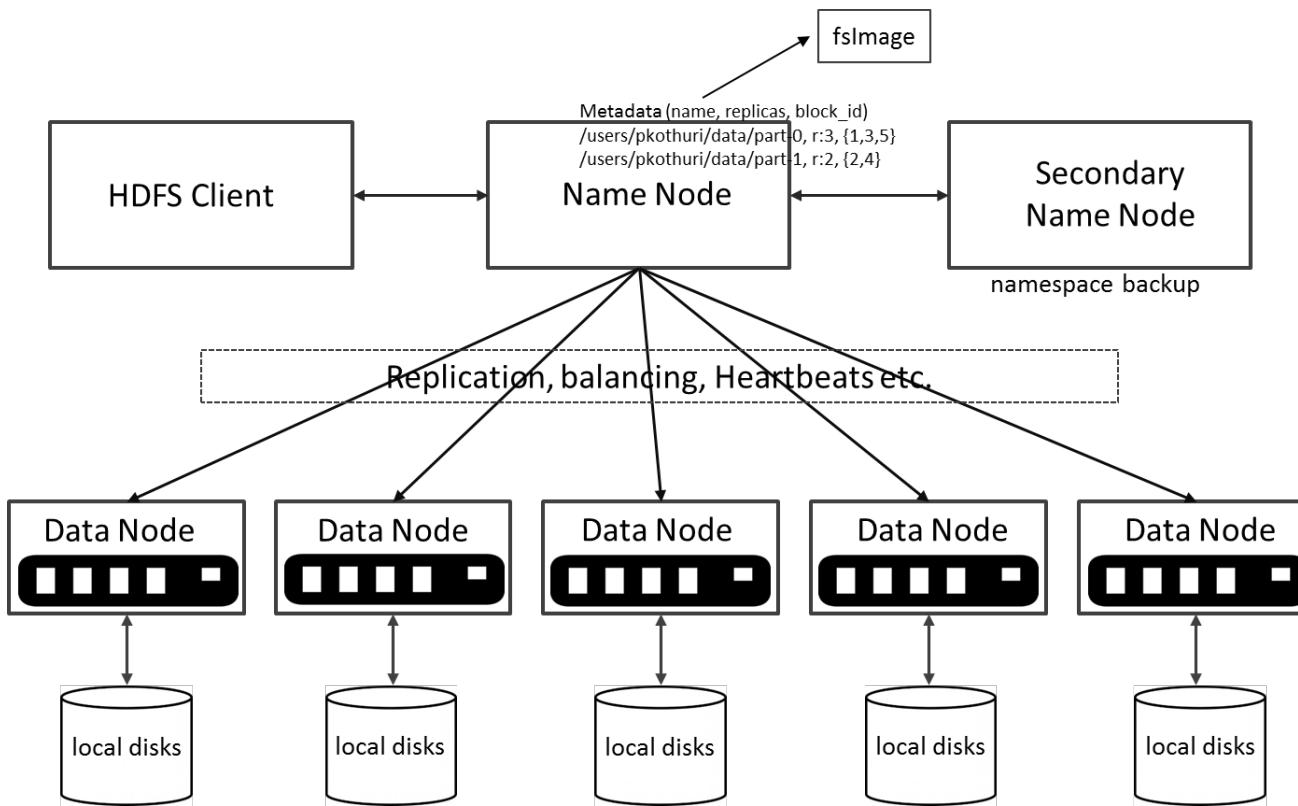
# HDFS – Data Organization

- Each file written into HDFS is split into data blocks
- Each block is stored on one or more nodes
- Each copy of the block is called replica
- Block placement policy
  - First replica is placed on the local node
  - Second replica is placed in a different rack
  - Third replica is placed in the same rack as the second replica

# HDFS architecture (1)



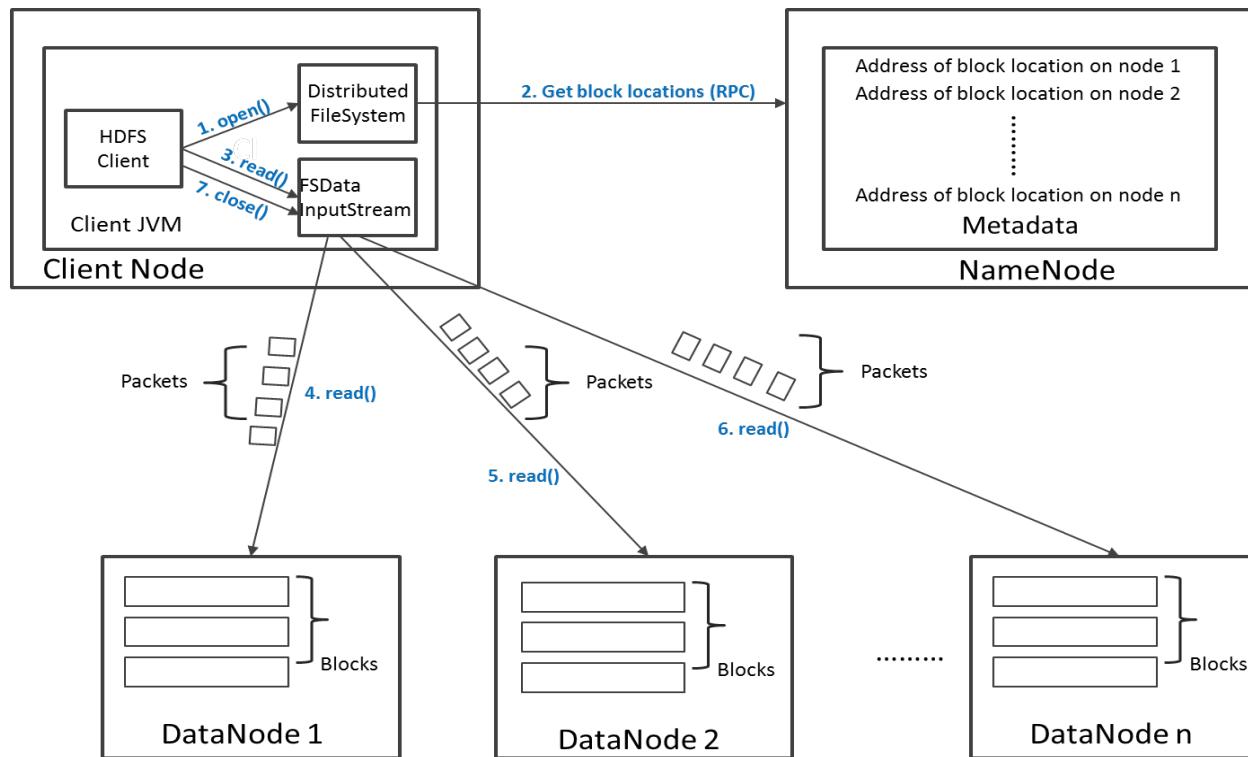
# HDFS architecture (2)



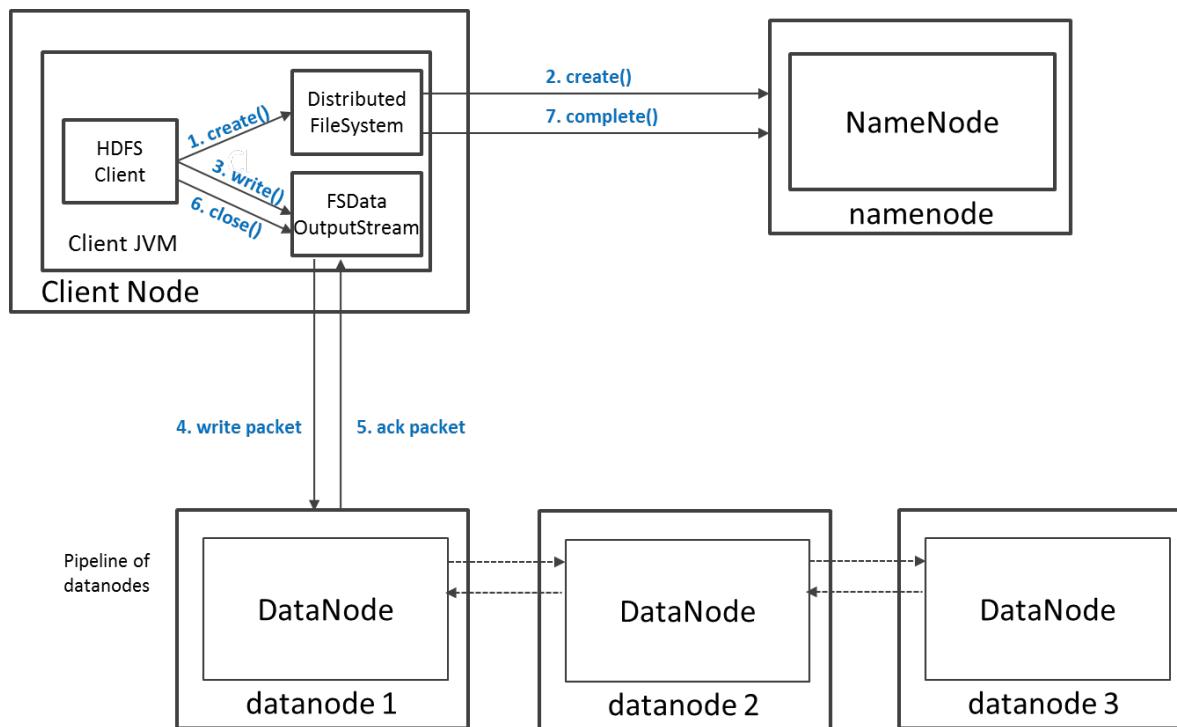
There are two (*and a half*) types of machines in a HDFS cluster

- **NameNode** is the heart of an HDFS filesystem, it maintains and manages the file system metadata. E.g; what blocks make up a file, and on which datanodes those blocks are stored
- **DataNode** where HDFS stores the actual data, there are usually quite a few of these

# Read operation in HDFS



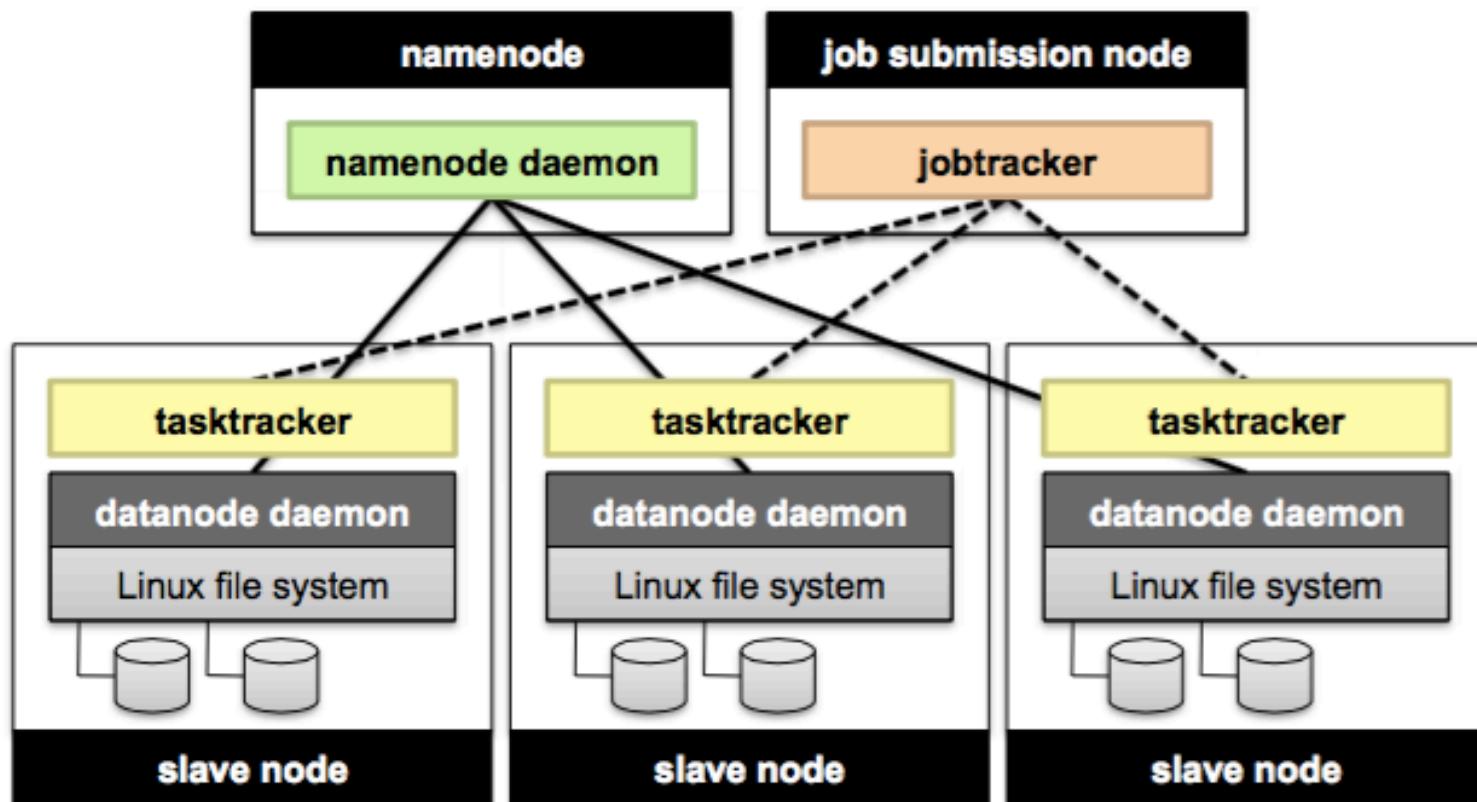
# Write operation in HDFS



# Unique features of HDFS

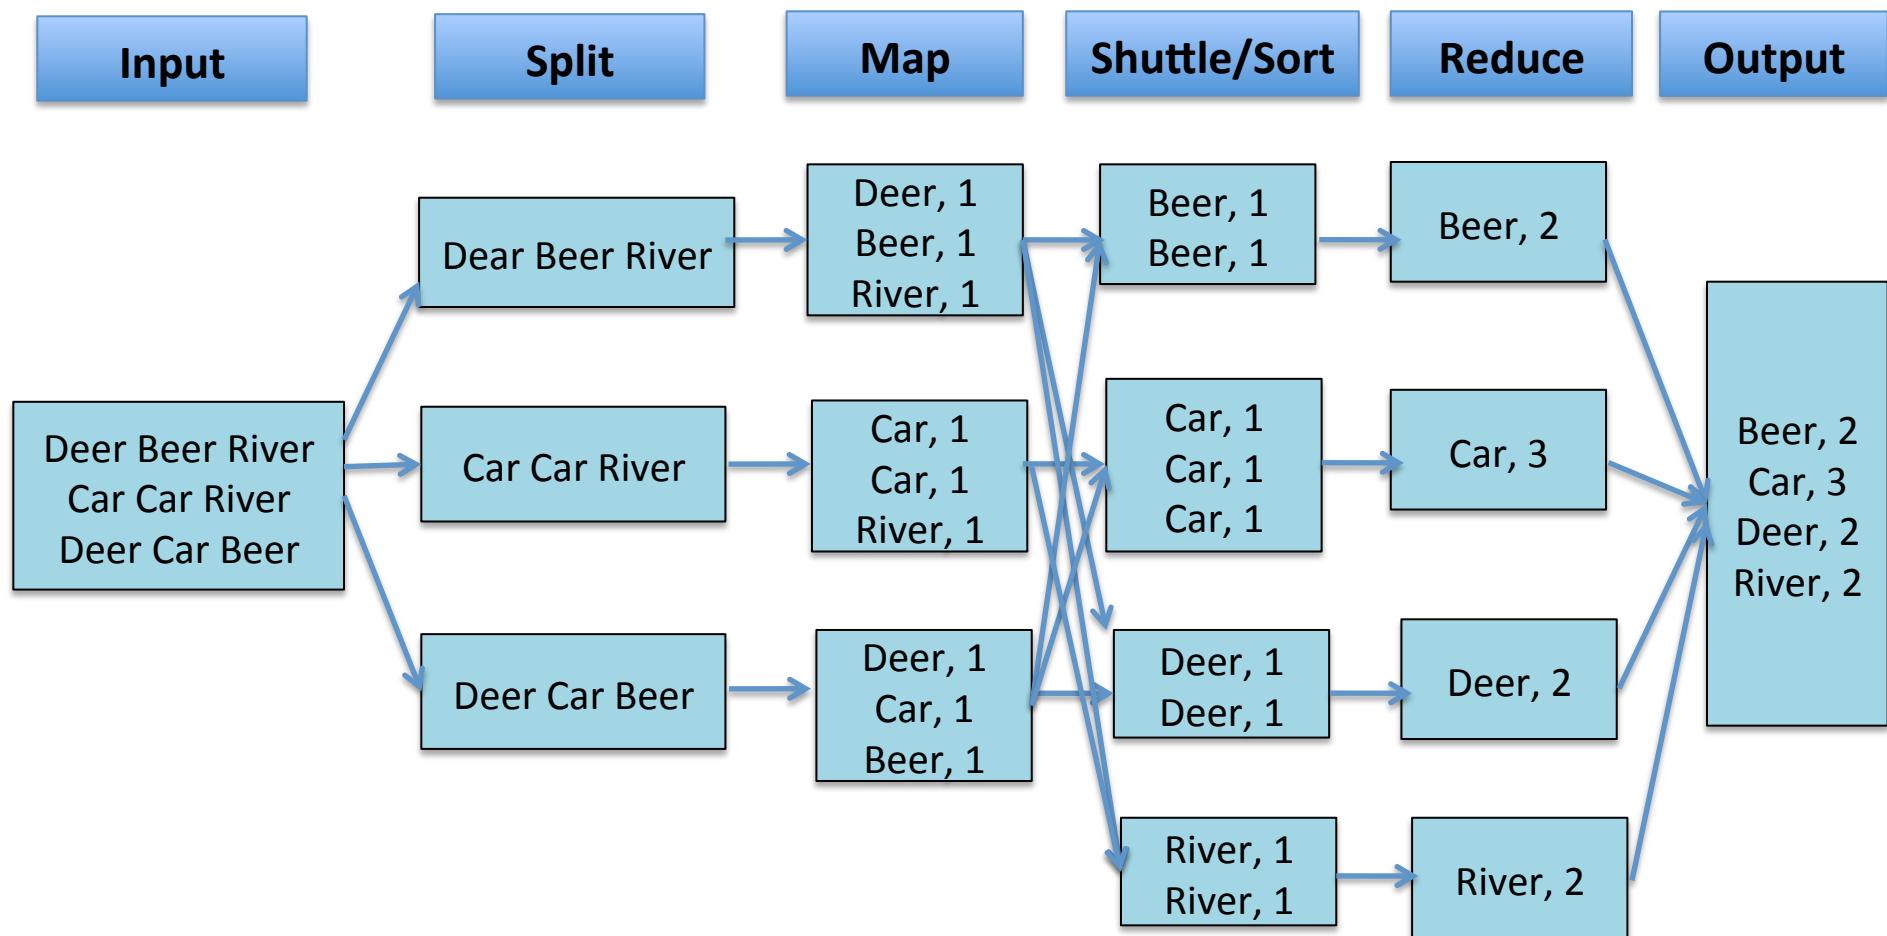
- HDFS also has a bunch of unique features that make it ideal for distributed systems:
  - **Failure tolerant** - data is duplicated across multiple DataNodes to protect against machine failures. The default is a replication factor of 3 (every block is stored on three machines).
  - **Scalability** - data transfers happen directly with the DataNodes so your read/write capacity scales fairly well with the number of DataNodes
  - **Space** - need more disk space? Just add more DataNodes and re-balance
  - **Industry standard** - Other distributed applications are built on top of HDFS (HBase, Map-Reduce)
- HDFS is designed to process large data sets with **write-once-read-many**, **it is not for low latency access**

# MapReduce & HDFS

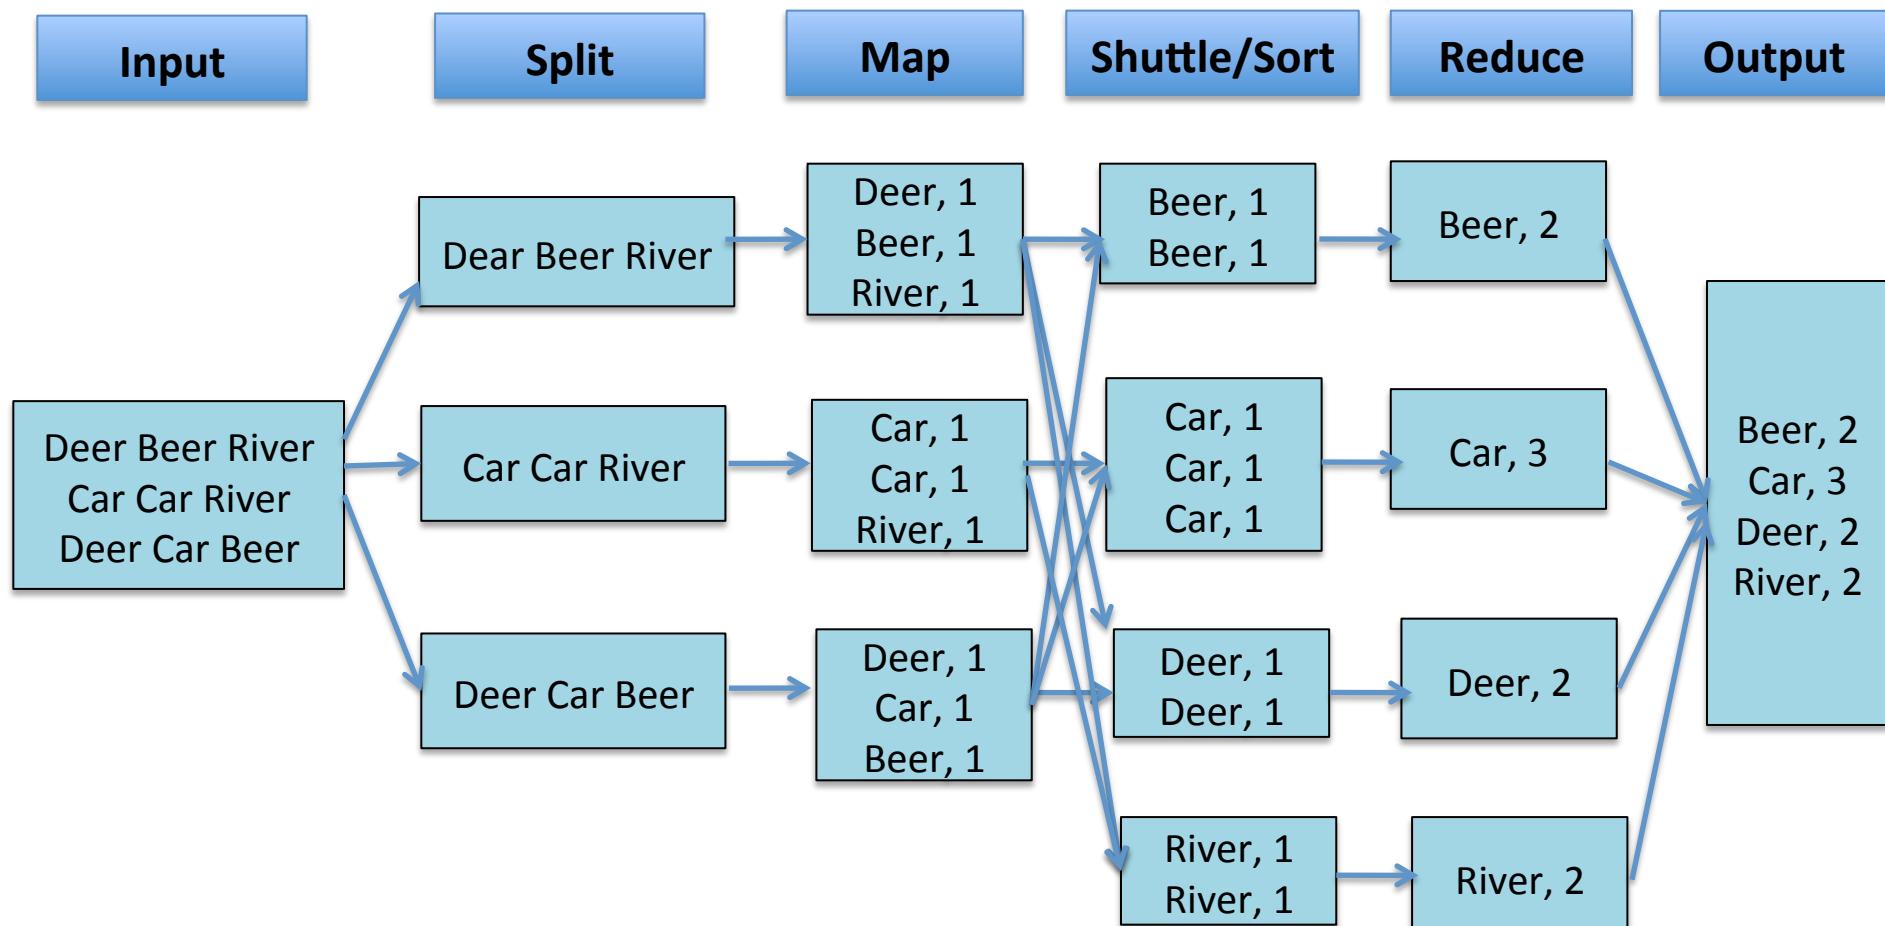


# Algorithm & programming

# MapReduce Example: Word Count



# MapReduce Example: Word Count



**Q: What are the Key and Value Pairs of Map and Reduce?**

**Map:** Key=word, Value=1

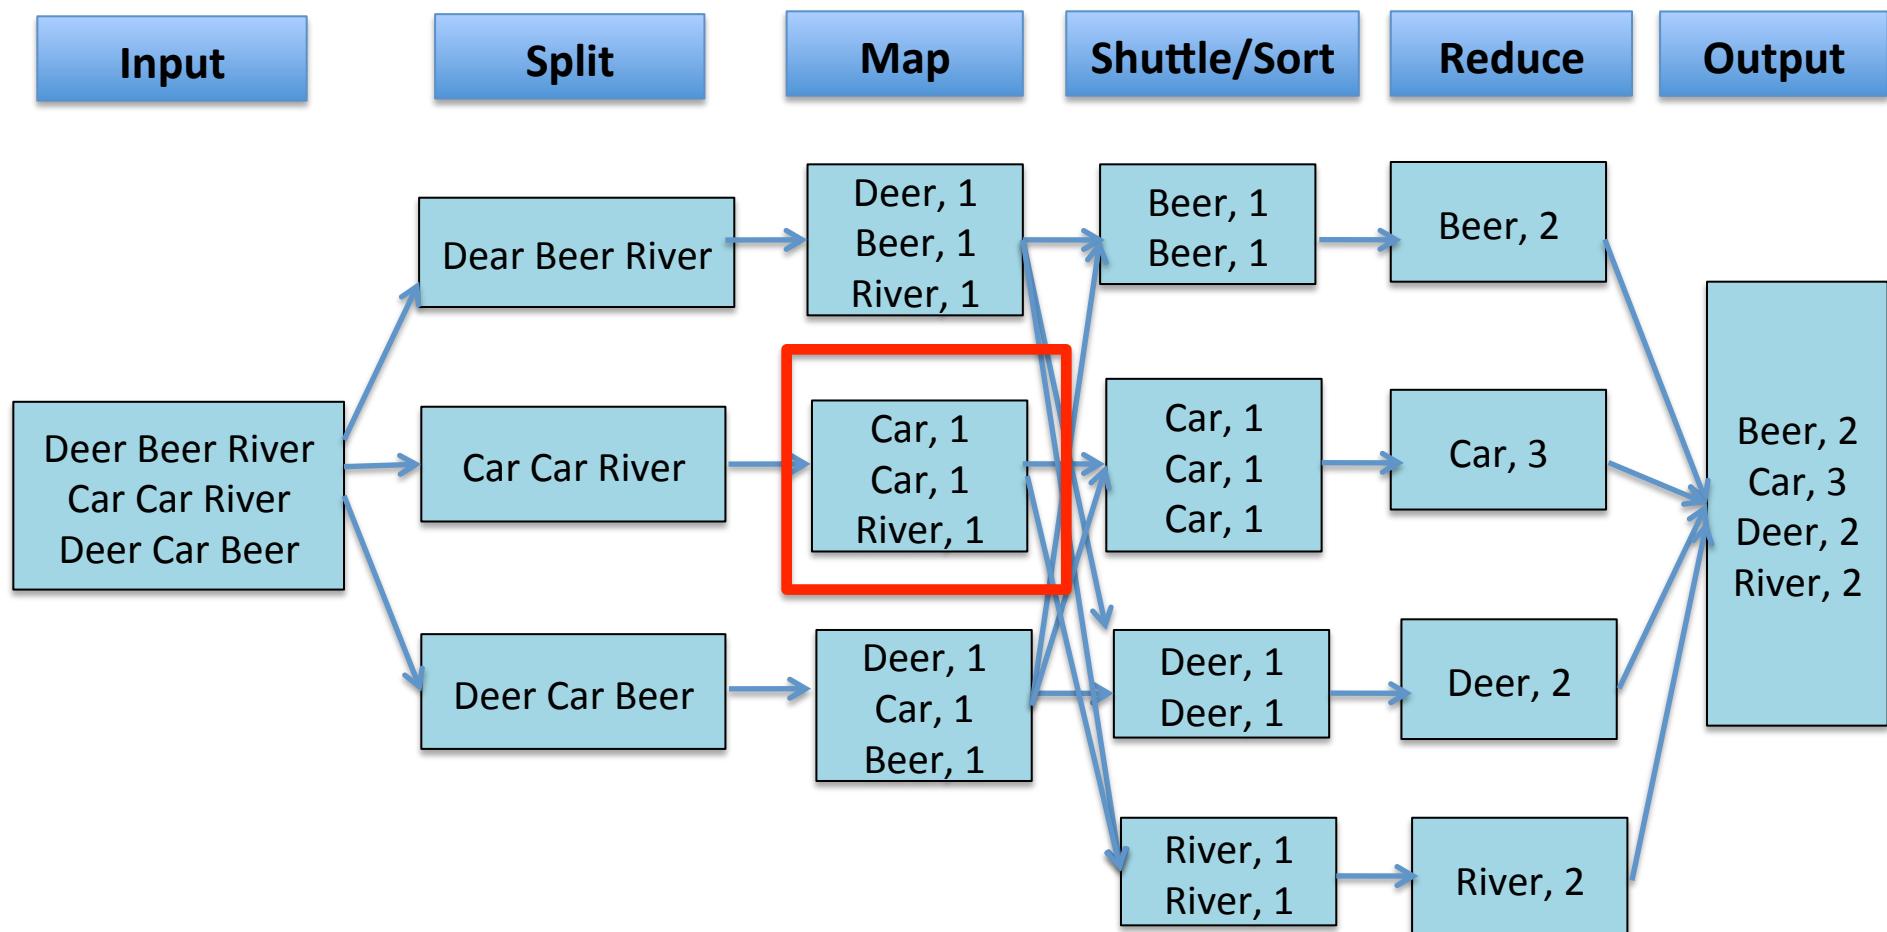
**Reduce:** Key=word, Value=aggregated count

# Word Count: baseline

```
1: class MAPPER
2:   method MAP(docid a, doc d)
3:     for all term t  $\in$  doc d do
4:       EMIT(term t, count 1)

1: class REDUCER
2:   method REDUCE(term t, counts [c1, c2, ...])
3:     sum  $\leftarrow$  0
4:     for all count c  $\in$  counts [c1, c2, ...] do
5:       sum  $\leftarrow$  sum + c
6:     EMIT(term t, count s)
```

# MapReduce Example: Word Count

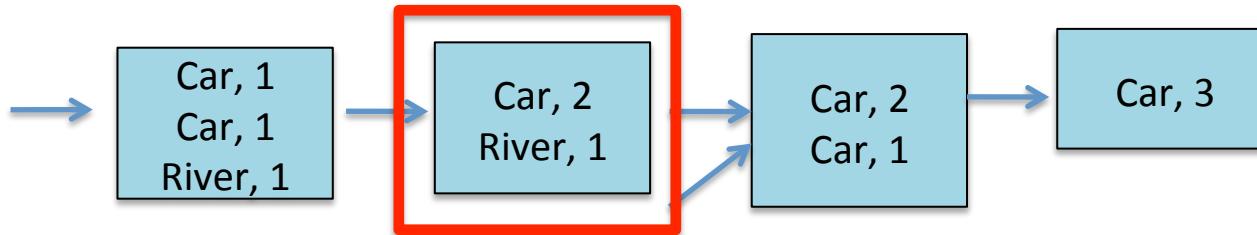


Q: Do you see any place we can improve the efficiency?

Local aggregation at mapper will be able to improve  
MapReduce efficiency.

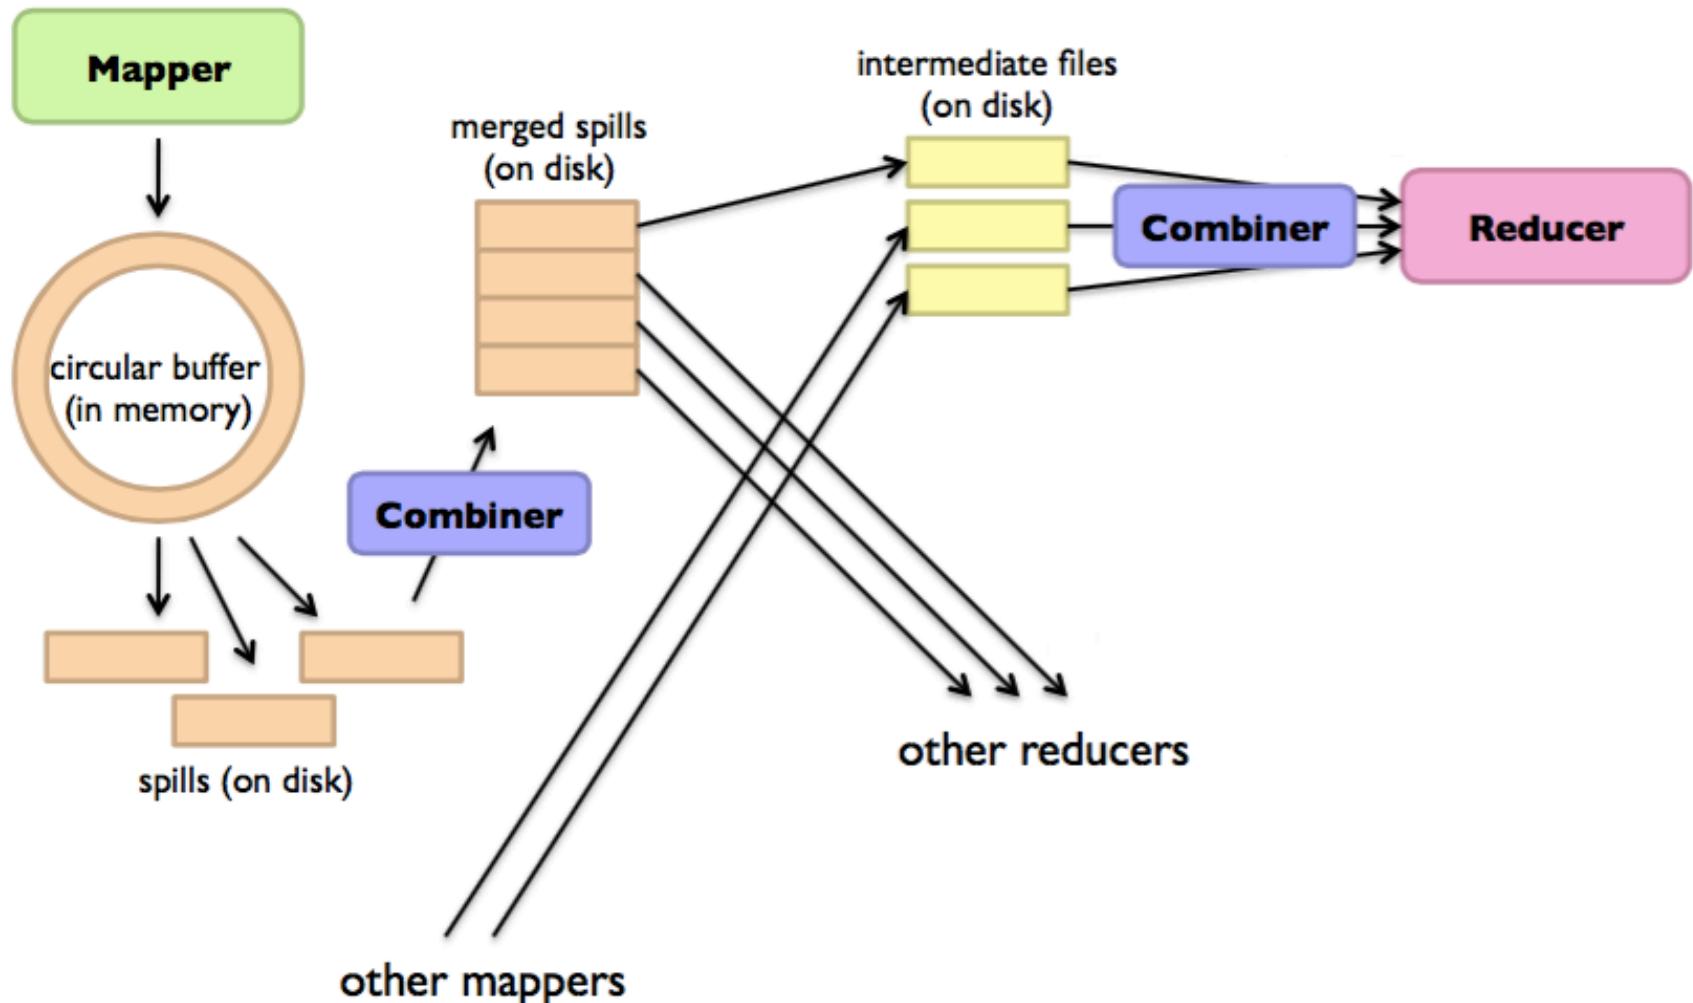
# MapReduce: Combiner

- **Combiner:** do local aggregation/combine task at mapper

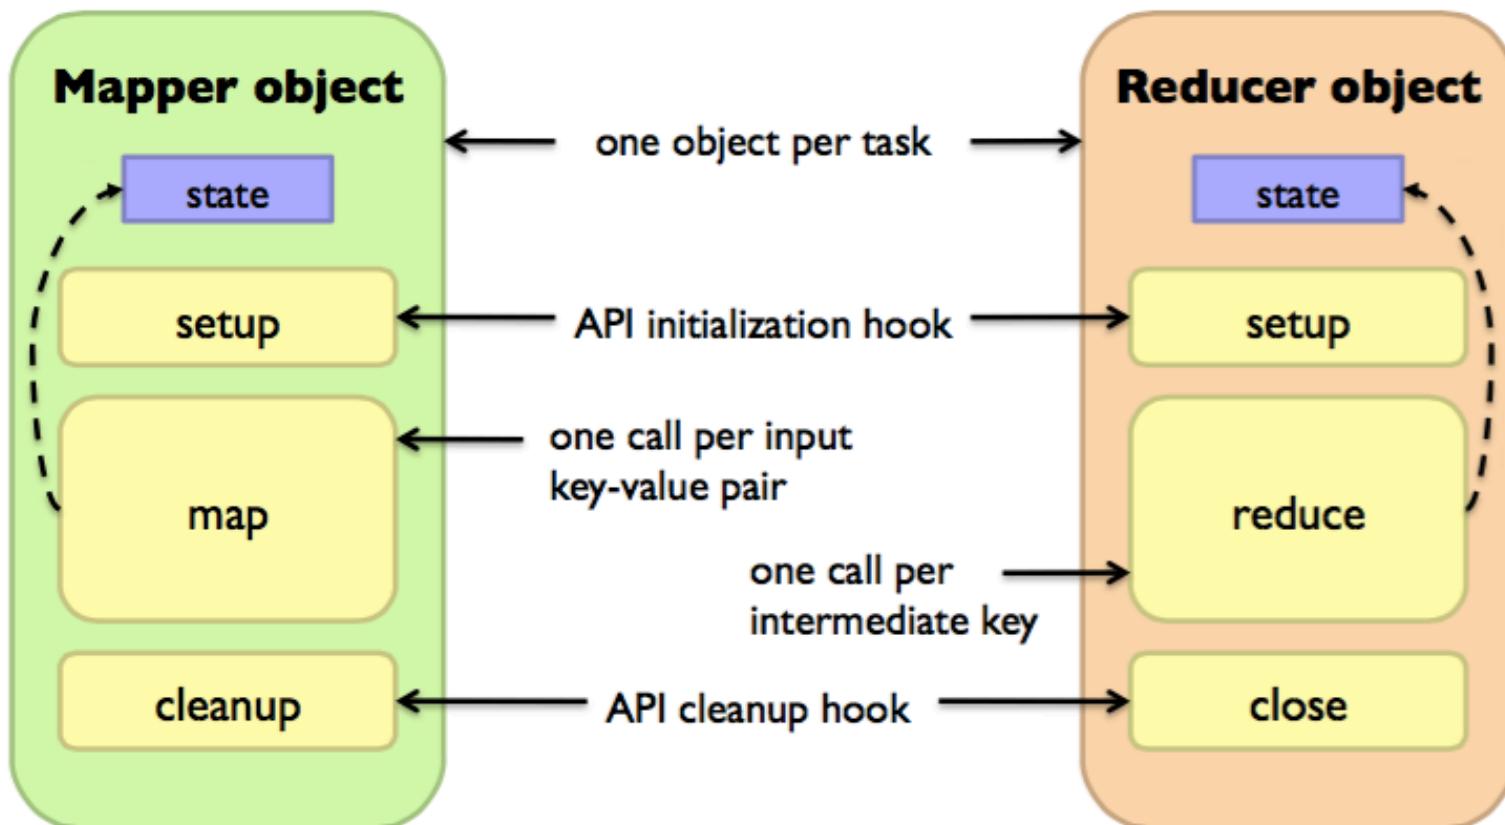


- **Q: What are the benefits of using combiner:**
  - Reduce memory/disk requirement of Map tasks
  - Reduce network traffic
- **Q: Can we remove the reduce function?**
  - No, reducer still needs to process records with same key but from *different mappers*
- **Q: How would you implement combiner?**
  - It is the same as Reducer!

# Shuffle and sort



# Preserving state



# Implementation don't

- Don't unnecessarily create objects
  - Object creation is costly
  - Garbage collection is costly
- Don't buffer objects
  - Processes have limited heap size (remember, commodity machines)
  - May work for small datasets, but won't scale!

# Word Count: version 1

```
1: class MAPPER
2:   method MAP(docid a, doc d)
3:     H  $\leftarrow$  new ASSOCIATIVEARRAY
4:     for all term t  $\in$  doc d do
5:       H{t}  $\leftarrow$  H{t} + 1            $\triangleright$  Tally counts for entire document
6:     for all term t  $\in$  H do
7:       EMIT(term t, count H{t})
```

# Word Count: version 2

```
1: class MAPPER
2:   method INITIALIZE
3:      $H \leftarrow$  new ASSOCIATIVEARRAY
4:   method MAP(docid  $a$ , doc  $d$ )
5:     for all term  $t \in$  doc  $d$  do
6:        $H\{t\} \leftarrow H\{t\} + 1$ 
7:   method CLOSE
8:     for all term  $t \in H$  do
9:       EMIT(term  $t$ , count  $H\{t\}$ )
```

Key idea: preserve state across  
input key-value pairs!

▷ Tally counts *across* documents

Are combiners still need?

# Design pattern for local aggregation

- “In-mapper combining”
  - Fold the functionality of the combiner into the mapper by preserving state across multiple map calls
- Advantages
  - Speed
  - [Why is this faster than actual combiners?](#)
- Disadvantages
  - Explicit memory management required
  - Potential for order-dependent bugs

# Combiner design

- Combiners and reducers share same method signature
  - Sometimes, reducers can serve as combiners
  - Often, not...
- Remember: combiner are optional optimizations
  - Should not affect algorithm correctness
  - May be run 0, 1, or multiple times
- Example: find average of integers associated with the same key

# Computing the Mean: version 1

```
1: class MAPPER
2:     method MAP(string t, integer r)
3:         EMIT(string t, integer r)

1: class REDUCER
2:     method REDUCE(string t, integers [r1, r2, ...])
3:         sum ← 0
4:         cnt ← 0
5:         for all integer r ∈ integers [r1, r2, ...] do
6:             sum ← sum + r
7:             cnt ← cnt + 1
8:             ravg ← sum/cnt
9:             EMIT(string t, integer ravg)
```

Why can't we use Reducer as Combiner?

# Computing the Mean: version 2

```
1: class MAPPER
2:     method MAP(string t, integer r)
3:         EMIT(string t, integer r)
4:
5: class COMBINER
6:     method COMBINE(string t, integers [r1, r2, ...])
7:         sum ← 0
8:         cnt ← 0
9:         for all integer r ∈ integers [r1, r2, ...] do
10:             sum ← sum + r
11:             cnt ← cnt + 1
12:         EMIT(string t, pair (sum, cnt))           ▷ Separate sum and count
13:
14: class REDUCER
15:     method REDUCE(string t, pairs [(s1, c1), (s2, c2) ...])
16:         sum ← 0
17:         cnt ← 0
18:         for all pair (s, c) ∈ pairs [(s1, c1), (s2, c2) ...] do
19:             sum ← sum + s
20:             cnt ← cnt + c
21:         ravg ← sum/cnt
22:         EMIT(string t, integer ravg)
```

**Why doesn't this work?**

# Computing the Mean: version 3

```
1: class MAPPER
2:     method MAP(string t, integer r)
3:         EMIT(string t, pair (r, 1))
4:
5: class COMBINER
6:     method COMBINE(string t, pairs [(s1, c1), (s2, c2) ...])
7:         sum ← 0
8:         cnt ← 0
9:         for all pair (s, c) ∈ pairs [(s1, c1), (s2, c2) ...] do
10:             sum ← sum + s
11:             cnt ← cnt + c
12:         EMIT(string t, pair (sum, cnt))
13:
14: class REDUCER
15:     method REDUCE(string t, pairs [(s1, c1), (s2, c2) ...])
16:         sum ← 0
17:         cnt ← 0
18:         for all pair (s, c) ∈ pairs [(s1, c1), (s2, c2) ...] do
19:             sum ← sum + s
20:             cnt ← cnt + c
21:         ravg ← sum/cnt
22:         EMIT(string t, pair (ravg, cnt))
```

# Computing the Mean: version 4

```
1: class MAPPER
2:     method INITIALIZE
3:          $S \leftarrow$  new ASSOCIATIVEARRAY
4:          $C \leftarrow$  new ASSOCIATIVEARRAY
5:     method MAP(string  $t$ , integer  $r$ )
6:          $S\{t\} \leftarrow S\{t\} + r$ 
7:          $C\{t\} \leftarrow C\{t\} + 1$ 
8:     method CLOSE
9:         for all term  $t \in S$  do
10:             EMIT(term  $t$ , pair ( $S\{t\}$ ,  $C\{t\}$ ))
```

Are combiners still need?

# Word Count & sorting

- **New Goal:** output all words sorted by their frequencies (total counts) in a document.
- **Question:** How would you adopt the basic word count program to solve it?
- **Solution:**
  - Sort words by their counts in the reducer
  - Problem: what happens if we have more than one reducer?

# Word Count & sorting

- **New Goal:** output all words sorted by their frequencies (total counts) in a document.
- **Question:** How would you adopt the basic word count program to solve it?
- **Solution:**
  - Do two rounds of MapReduce
  - In the 2<sup>nd</sup> round, take the output of WordCount as input but switch key and value pair!
  - Leverage the sorting capability of *shuffle/sort* to do the global sorting!

# Word Count & top K words

- **New Goal:** output the top K words sorted by their frequencies (total counts) in a document.
- **Question:** How would you adopt the basic word count program to solve it?
- **Solution:**
  - Use the solution of previous problem and only grab the top K in the final output
  - Problem: is there a more efficient way to do it?

# Word Count & top K words

- **New Goal:** output the top K words sorted by their frequencies (total counts) in a document.
- **Question:** How would you adopt the basic word count program to solve it?
- **Solution:**
  - Add a sort function to the *reducer* in the first round and only output the top K words
  - Intuition: the global top K must be a local top K in any reducer!

# MapReduce In-class Exercise

- **Problem:** Find the maximum monthly temperature for each year from weather reports
- **Input:** A set of records with format as:  
 $\langle \text{Year/Month}, \text{Average Temperature of that month} \rangle$ 
  - **(200707,100), (200706,90)**
  - **(200508, 90), (200607,100)**
  - **(200708, 80), (200606,80)**
- **Question:** write down the Map and Reduce function to solve this problem
  - Assume we split the input by line

# Mapper and Reducer of Max Temperature

- Map(key, value){  
    // key: line number  
    // value: tuples in a line  
    for each tuple t in value:  
        Emit(t->year, t->temperature);}
- Reduce(key, list of values){  
    // key: year  
    //list of values: a list of monthly temperature  
    int max\_temp = -100;  
    for each v in values:  
        max\_temp= max(v, max\_temp);  
    Emit(key, max\_temp);}

Combiner is the same as Reducer

# MapReduce Example: Max Temperature

Input

(200707,100), (200706,90)  
(200508, 90), (200607,100)  
(200708, 80), (200606,80)

Map

(2007,100), (2007,90)

(2005, 90), (2006,100)

(2007, 80), (2006, 80)

Combine

**(2007,100)**

(2005, 90), (2006,100)

(2007, 80), (2006, 80)

Shuttle/Sort

(2005,[90])

(2006,[100, 80])

(2007,[100, 80])

Reduce

**(2005,90)**

**(2006,100)**  
HPC Lab-CSE-HCMUT

**(2007,100)**

# MapReduce In-class Exercise

- **Key-Value Pair of Map and Reduce:**
  - **Map:** (year, temperature)
  - **Reduce:** (year, maximum temperature of the year)
- **Question: How to use the above Map Reduce program (*that contains the combiner*) with slight changes to find the average monthly temperature of the year?**

# Mapper and Reducer of Average Temperature

- Map(key, value){  
    // key: line number  
    // value: tuples in a line  
    for each tuple t in value:  
        Emit(t->year, t->temperature);}
  - Reduce(key, list of values){  
    // key: year  
    // list of values: a list of monthly temperatures  
    int total\_temp = 0;  
    for each v in values:  
        total\_temp= total\_temp+v;  
    Emit(key, total\_temp/size\_of(values));}
- Combiner is the same as Reducer

# MapReduce Example: Average Temperature

Input

(200707,100), (200706,90)  
(200508, 90), (200607,100)  
(200708, 80), (200606,80)

Real average of  
2007: 90

Map

(2007,100), (2007,90)

(2005, 90), (2006,100)

(2007, 80), (2006,80)

Combine

(2007,95)

(2005, 90), (2006,100)

(2007, 80), (2006,80)

Shuttle/Sort

(2005,[90])

(2006,[100, 80])

(2007,[95, 80])

Reduce

(2005,90)

(2006,90)  
HPC Lab-CSE-HCMUT

(2007,87.5)

# MapReduce In-class Exercise

- The problem is with the combiner!
- Here is a simple counterexample:
  - $(2007, 100), (2007, 90) \rightarrow (2007, 95)$   
 $(2007, 80) \rightarrow (2007, 80)$
  - Average of the above is:  $(2007, 87.5)$
  - However, the real average is:  $(2007, 90)$
- However, we can do a small trick to get around this
  - Mapper:  $(2007, 100), (2007, 90) \rightarrow (2007, <190, 2>)$   
 $(2007, 80) \rightarrow (2007, <80, 1>)$
  - Reducer:  $(2007, <270, 3>) \rightarrow (2007, 90)$

# MapReduce Example: Average Temperature

Input

(200707,100), (200706,90)  
(200508, 90), (200607,100)  
(200708, 80), (200606,80)

Map

(2007,100), (2007,90)

(2005, 90), (2006,100)

(2007, 80), (2006,80)

Combine

(2007,<190,2>)

(2005, <90,1>),  
(2006, <100,1>)

(2007, <80,1>),  
(2006,<80,1>)

Shuttle/Sort

(2005,[<90,1>])

(2006,[<100,1>, <80,1>])

(2007,[<190,2>, <80,1>])

Reduce

(2005,90)

(2006,90)  
HPC Lab-CSE-HCMUT

(2007,90)

# Mapper and Reducer of Average Temperature

- **Map(key, value){**  
    // key: line number  
    // value: tuples in a line  
    for each tuple t in value:  
        Emit(t->year, t->temperature);}
- **Reduce (key, list of values){**  
    // key: year  
    // list of values: a list of <temperature sums, counts> tuples  
  
    int total\_temp = 0;  
    int total\_count=0;  
    for each v in values:  
        total\_temp= total\_temp+v->sum;  
        total\_count=total\_count+v->count;  
  
    **Emit(key,**total\_temp/total\_count**)**
- **Combine(key, list of values){**  
    // key: year  
    // list of values: a list of monthly temperature  
  
    int total\_temp = 0;  
    for each v in values:  
        total\_temp= total\_temp+v;  
  
    **Emit(key,<total\_temp,size\_of(values)>);**

# MapReduce In-class Exercise

- Functions that can use combiner are called *distributive*:
  - Distributive: Min/Max(), Sum(), Count(), TopK()
  - Non-distributive: Mean(), Median(), Rank()

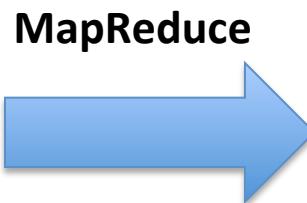
Gray, Jim\*, et al. "Data cube: A relational aggregation operator generalizing group-by, cross-tab, and sub-totals." Data Mining and Knowledge Discovery 1.1 (1997): 29-53.

\*Jim Gray received Turing Award in 1998

# Map Reduce Problems Discussion

- **Problem 1:** Find Word Length Distribution
- **Statement:** Given a set of documents, use MapReduce to find the length distribution of all words contained in the documents
- **Question:**
  - What are the Mapper and Reducer Functions?

This is a test data for  
the word length  
distribution problem



12: 1  
7: 1  
6: 1  
4: 4  
3: 2  
2: 1  
1: 1

# Mapper and Reducer of Word Length Distribution

- Map(key, value){  
    // key: document name  
    // value: words in a document  
    for each word w in value:  
        Emit(length(w), w);}
- Reduce(key, list of values){  
    // key: length of a word  
    // list of values: a list of words with the same length  
    Emit(key, size\_of(values));}

# Map Reduce Problems Discussion

- **Problem 1:** Find Word Length Distribution
- **Mapper and Reducer:**
  - Mapper(document)  
  { Emit (**Length(word)**, **word**) }
  - Reducer(output of map)  
  { Emit (**Length(word)**, **Size of (List of words at a particular length)**)}}

# Map Reduce Problems Discussion

- **Problem 2:** Indexing & Page Rank
- **Statement:** Given a set of web pages, each page has a **page rank** associated with it, use Map-Reduce to find, for each word, a list of pages (sorted by rank) that contains that word
- **Question:**
  - What are the Mapper and Reducer Functions?

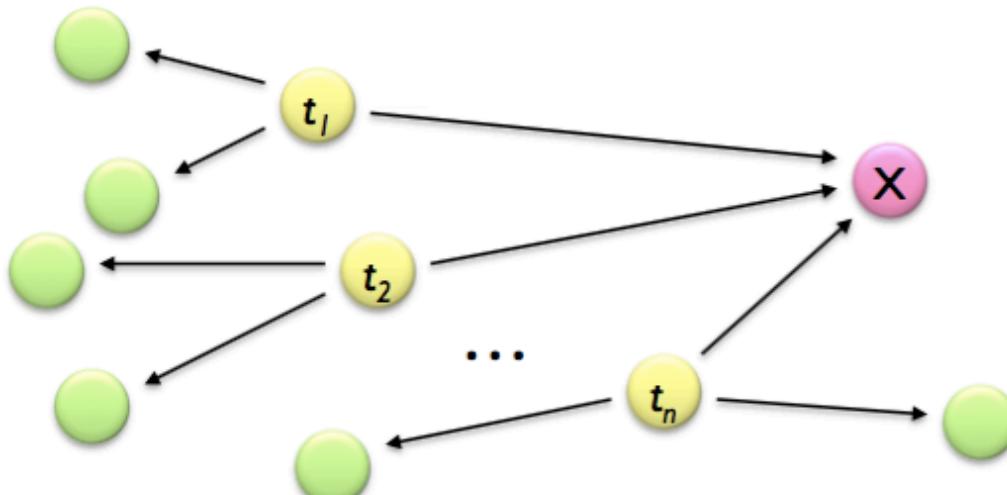


# Page Rank

Given page  $x$  with inlinks  $t_1, \dots, t_n$ , where

- $C(t)$  is the out-degree of  $t$
- $\alpha$  is probability of random jump
- $N$  is the total number of nodes in the graph

$$PR(x) = \alpha \left( \frac{1}{N} \right) + (1 - \alpha) \sum_{i=1}^n \frac{PR(t_i)}{C(t_i)}$$



# Mapper and Reducer of Indexing and PageRank

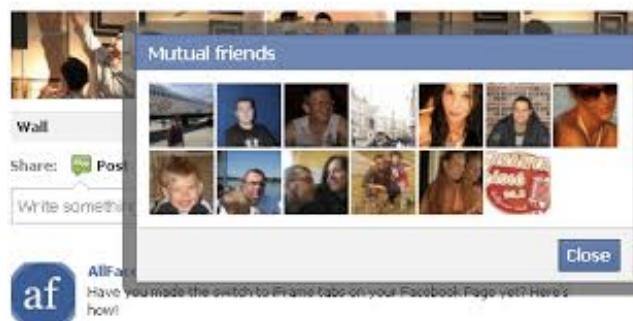
- Map(key, value){  
    // key: a page  
    // value: words in a page  
    for each word w in value:  
        Emit(w, <page\_id, page\_rank>);}
- Reduce(key, list of values){  
    // key: a word  
    // list of values: a list of pages containing that word  
    sorted\_pages=sort(values, page\_rank)  
    Emit(key, sorted\_pages);}

# Map Reduce Problems Discussion

- Problem 2: Indexing and Page Rank
- **Mapper and Reducer:**
  - Mapper(page\_id, <page\_text, page\_rank>)  
  { Emit (word, <page\_id, page\_rank>) }
  - Reducer(output of map)  
  { Emit (word, List of pages contains the word sorted by their page\_ranks)}

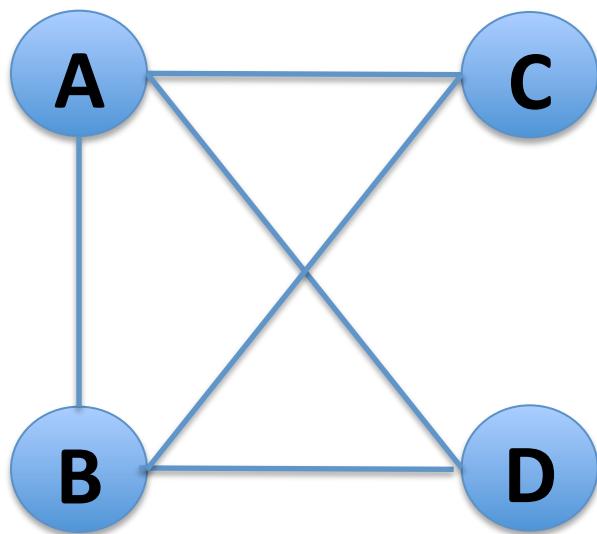
# Map Reduce Problems Discussion

- **Problem 3: Find Common Friends**
- **Statement:** Given a group of people on online social media (e.g., Facebook), each has a list of friends, use Map-Reduce to find common friends of any two persons who are friends
- **Question:**
  - What are the Mapper and Reducer Functions?



# Map Reduce Problems Discussion

- Problem 3: Find Common Friends
- Simple example:



Input:

A -> B,C,D

B-> A,C,D

C-> A,B

D->A,B

MapReduce

Output:

(A ,B) -> C,D

(A,C) -> B

(A,D) -> ..

....

# Mapper and Reducer of Common Friends

- Map(key, value){  
    // key: person\_id  
    // value: the list of friends of the person  
    for each friend f\_id in value:  
        Emit(<person\_id, f\_id>, value);}
- Reduce(key, list of values){  
    // key: <friend pair>  
    // list of values: a set of friend lists related with the friend pair  
    for v1, v2 in values:  
        common\_friends = v1 intersects v2;  
    Emit(key, common\_friends);}

# Map Reduce Problems Discussion

- **Problem 3:** Find Common Friends
- **Mapper and Reducer:**
  - Mapper(friend list of a person)
    - { for each person in the friend list:  
    Emit (<friend pair>, <list of friends>) }
  - Reducer(output of map)
    - { Emit (<friend pair>, **Intersection of two (i.e, the one in friend pair) friend lists**) }

# Map Reduce Problems Discussion

- Problem 3: Find Common Friends
- Mapper and Reducer:

Input:

A -> B,C,D  
B-> A,C,D  
C-> A,B  
D->A,B

Map:

(A,B) -> B,C,D  
(A,C) -> B,C,D  
(A,D) -> B,C,D  
**(A,B) -> A,C,D**  
(B,C) -> A,C,D  
(B,D) -> A,C,D  
(A,C) -> A,B  
(B,C) -> A,B  
(A,D) -> A,B  
(B,D) -> A,B

Reduce:

**(A,B)** -> C,D  
(A,C) -> B  
(A,D) -> B  
(B,C) -> A  
(B,D) -> A

*Suggest  
Friends ☺*

*Enjoy MR and Hadoop* ☺

