

Calculating π in Parallel Using MPI

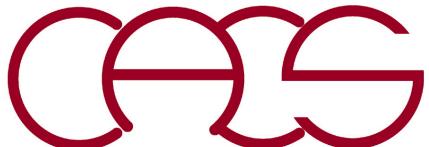
Aiichiro Nakano

*Collaboratory for Advanced Computing & Simulations
Department of Computer Science
Department of Physics & Astronomy
Department of Chemical Engineering & Materials Science
Department of Biological Sciences
University of Southern California*

Email: anakano@usc.edu

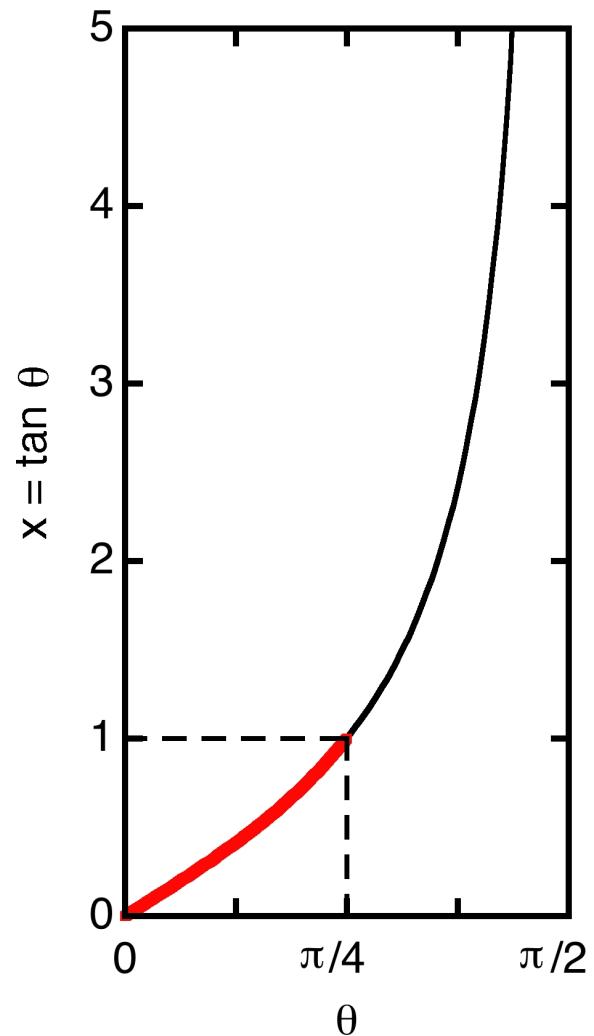
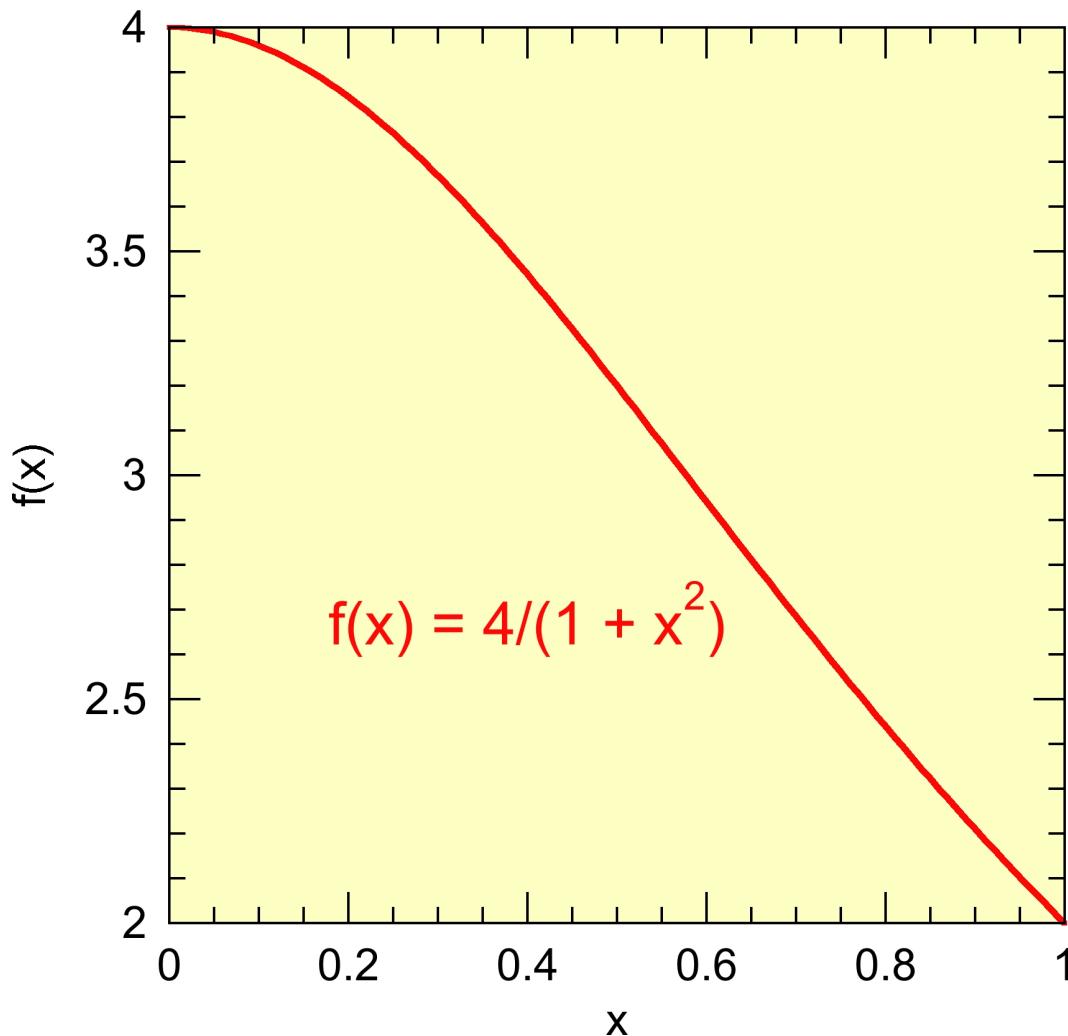
Objectives

1. Task decomposition (parallel programming = who does what)
2. Scalability analysis



Integral Representation of π

$$\int_0^1 dx \frac{4}{1+x^2} = \int_0^{\pi/4} \frac{d\theta}{\cos^2 \theta} \frac{4}{1+\tan^2 \theta} = \int_0^{\pi/4} 4d\theta = \pi$$



Numerical Integration of π

- Integration

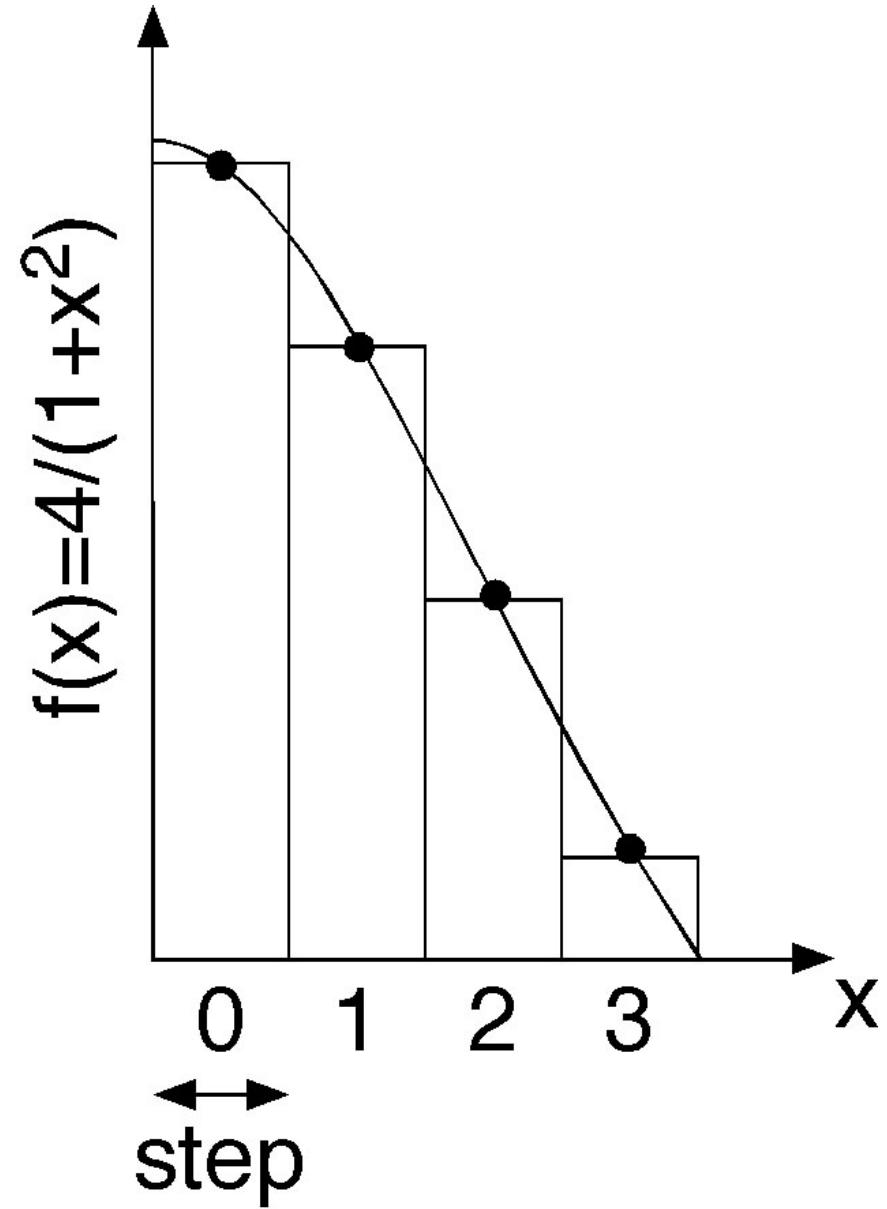
$$\int_0^1 dx \frac{4}{1+x^2} = \pi$$

- Discretization:

$$\Delta = 1/N: \text{step} = 1/\text{NBIN}$$
$$x_i = (i+0.5)\Delta \quad (i=0, \dots, N-1)$$

$$\sum_{i=0}^{N-1} \frac{4}{1+x_i^2} \Delta \cong \pi$$

```
#include <stdio.h>
#define NBIN 10000000
void main() {
    int i; double step,x,sum=0.0,pi;
    step = 1.0/NBIN;
    for (i=0; i<NBIN; i++) {
        x = (i+0.5)*step;
        sum += 4.0/(1.0+x*x);
    }
    pi = sum*step;
    printf( "PI = %f\n" ,pi);
}
```



Parallelization: Who Does What?

Interleaved assignment of quadrature points (bins) to MPI processes

Single program multiple data (SPMD)

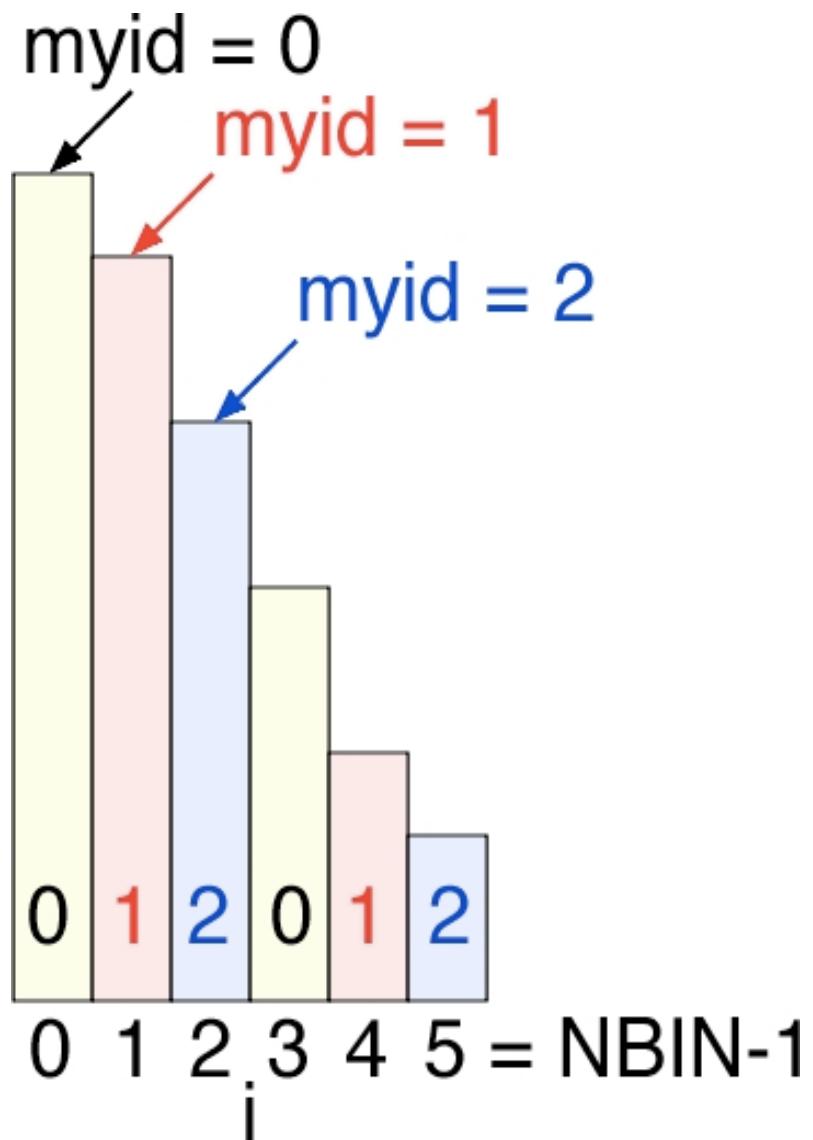
```
...
for (i=myid; i<NBIN; i+=nprocs)
{
    x = (i+0.5)*step;
    sum += 4.0/(1.0+x*x);
}
partial = sum*step;
pi = global_sum(partial);
...
```

`MPI_Comm_rank(MPI_COMM_WORLD, &myid)`

myid = MPI rank

nprocs = Number of MPI processes

`MPI_Comm_size(MPI_COMM_WORLD, &nprocs)`

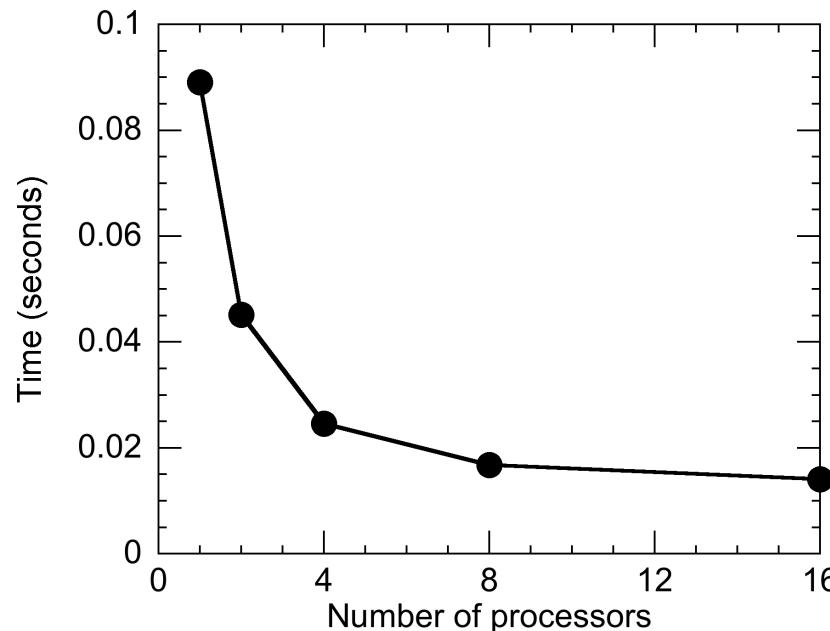


- Use `double MPI_Wtime()` to measure the running time in seconds

Parallel Running Time

global_pi.c: NBIN = 10⁷, on HPC

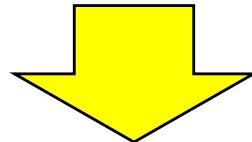
```
#SBATCH --ntasks-per-node=16
#SBATCH --nodes=1
...
srun -n $SLURM_NTASKS --mpi=pmi2 ./global_pi
srun -n 8 --mpi=pmi2 ./global_pi
srun -n 4 --mpi=pmi2 ./global_pi
srun -n 2 --mpi=pmi2 ./global_pi
srun -n 1 --mpi=pmi2 ./global_pi
```



How Efficient Is the Parallel Program?

Scalability Analysis

- Parallel computing = Solving a big problem (W) in a short time (T) using many processors (P)



- How W , T & P scale with each other?
- How to define the efficiency of a parallel program?

Parallel Efficiency

- Execution time: $T(W,P)$
 W : Workload
 P : Number of processors

- Speed: $S(W,P) = \frac{W}{T(W,P)}$

- Speedup: $S_P = \frac{S(W_P,P)}{S(W_1,1)} = \frac{W_P T(W_1,1)}{W_1 T(W_P,P)}$

- Efficiency: $E_P = \frac{S_P}{P} = \frac{W_P T(W_1,1)}{P W_1 T(W_P,P)}$

How to scale W_P with P ?

Fixed Problem-Size Scaling

$W_P = W - \text{constant}$ (strong scaling)

- **Speedup:** $S_P = \frac{T(W,1)}{T(W,P)}$
- **Efficiency:** $E_P = \frac{T(W,1)}{PT(W,P)}$

$$S_P = \frac{S(W_P, P)}{S(W_1, 1)} = \frac{W_P T(W_1, 1)}{W_1 T(W_P, P)}$$

$$E_P = \frac{S_P}{P} = \frac{W_P T(W_1, 1)}{P W_1 T(W_P, P)}$$

- **Amdahl's law:** f (= sequential fraction of the workload) limits the asymptotic speedup

$$\begin{aligned} T(W,P) &= fT(W,1) + \frac{(1-f)T(W,1)}{P} \\ \therefore S_P &= \frac{T(W,1)}{T(W,P)} = \frac{1}{f + (1-f)/P} \\ \therefore S_P &\rightarrow \frac{1}{f} \quad (P \rightarrow \infty) \end{aligned}$$

Isogranular Scaling

$W_P = Pw$ (weak scaling)

w = constant workload per processor (granularity)

- **Speedup:** $S_P = \frac{S(P \bullet w, P)}{S(w, 1)} = \frac{P \bullet w / T(P \bullet w, P)}{w / T(w, 1)} = \frac{P \bullet T(w, 1)}{T(P \bullet w, P)}$

- **Efficiency:** $E_P = \frac{S_P}{P} = \frac{T(w, 1)}{T(P \bullet w, P)}$

$$S_P = \frac{S(W_P, P)}{S(W_1, 1)} = \frac{W_P T(W_1, 1)}{W_1 T(W_P, P)}$$

$$E_P = \frac{S_P}{P} = \frac{W_P T(W_1, 1)}{P W_1 T(W_P, P)}$$

Analysis of Global_Pi Program

- Workload \propto Number of quadrature points, N (or `NBIN` in the program)
- Parallel execution time on P processors:
 - > Local computation $\propto N/P$

```
for (i=myid; i<N; i+=P){  
    x = (i+0.5)*step; partial += 4.0/(1.0+x*x);  
}
```

- > Butterfly computation/communication in `global()` $\propto \log P$

```
for (l=0; l<log2P; ++l) {  
    partner = myid XOR 2l;  
    send mydone to partner;  
    receive hisdone from partner;  
    mydone += hisdone  
}
```

$$\begin{aligned} T(N, P) &= T_{\text{comp}}(N, P) + T_{\text{global}}(P) \\ &= \alpha \frac{N}{P} + \beta \log P \end{aligned}$$

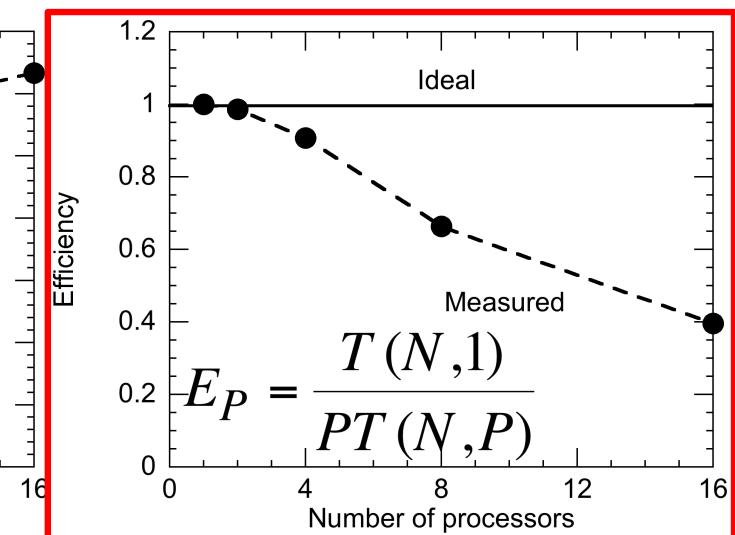
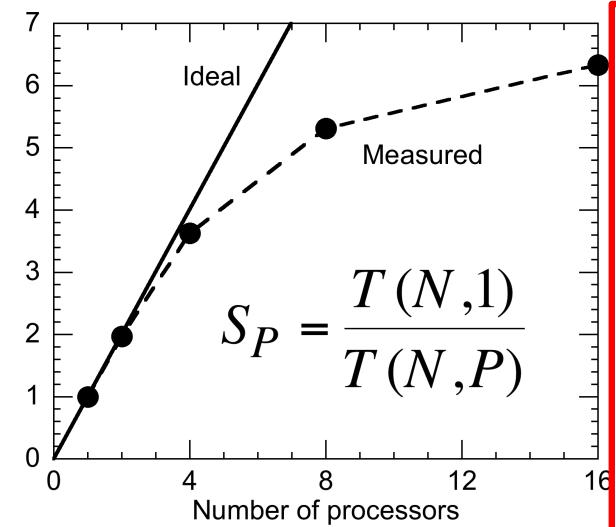
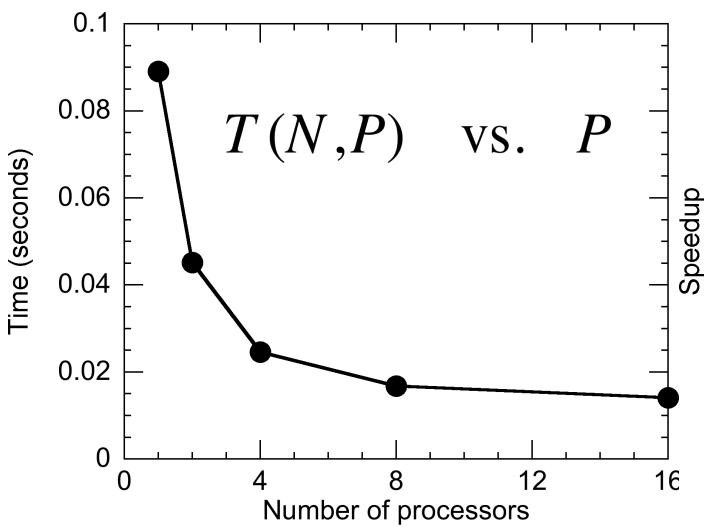
Fixed Problem-Size (Strong) Scaling

- Speedup:

$$S_P = \frac{T(N, 1)}{T(N, P)} = \frac{\alpha N}{\alpha \frac{N}{P} + \beta \log P} = \frac{P}{1 + \frac{\beta P \log P}{\alpha N}}$$

- Efficiency:

$$E_P = \frac{S_P}{P} = \frac{1}{1 + \frac{\beta P \log P}{\alpha N}}$$

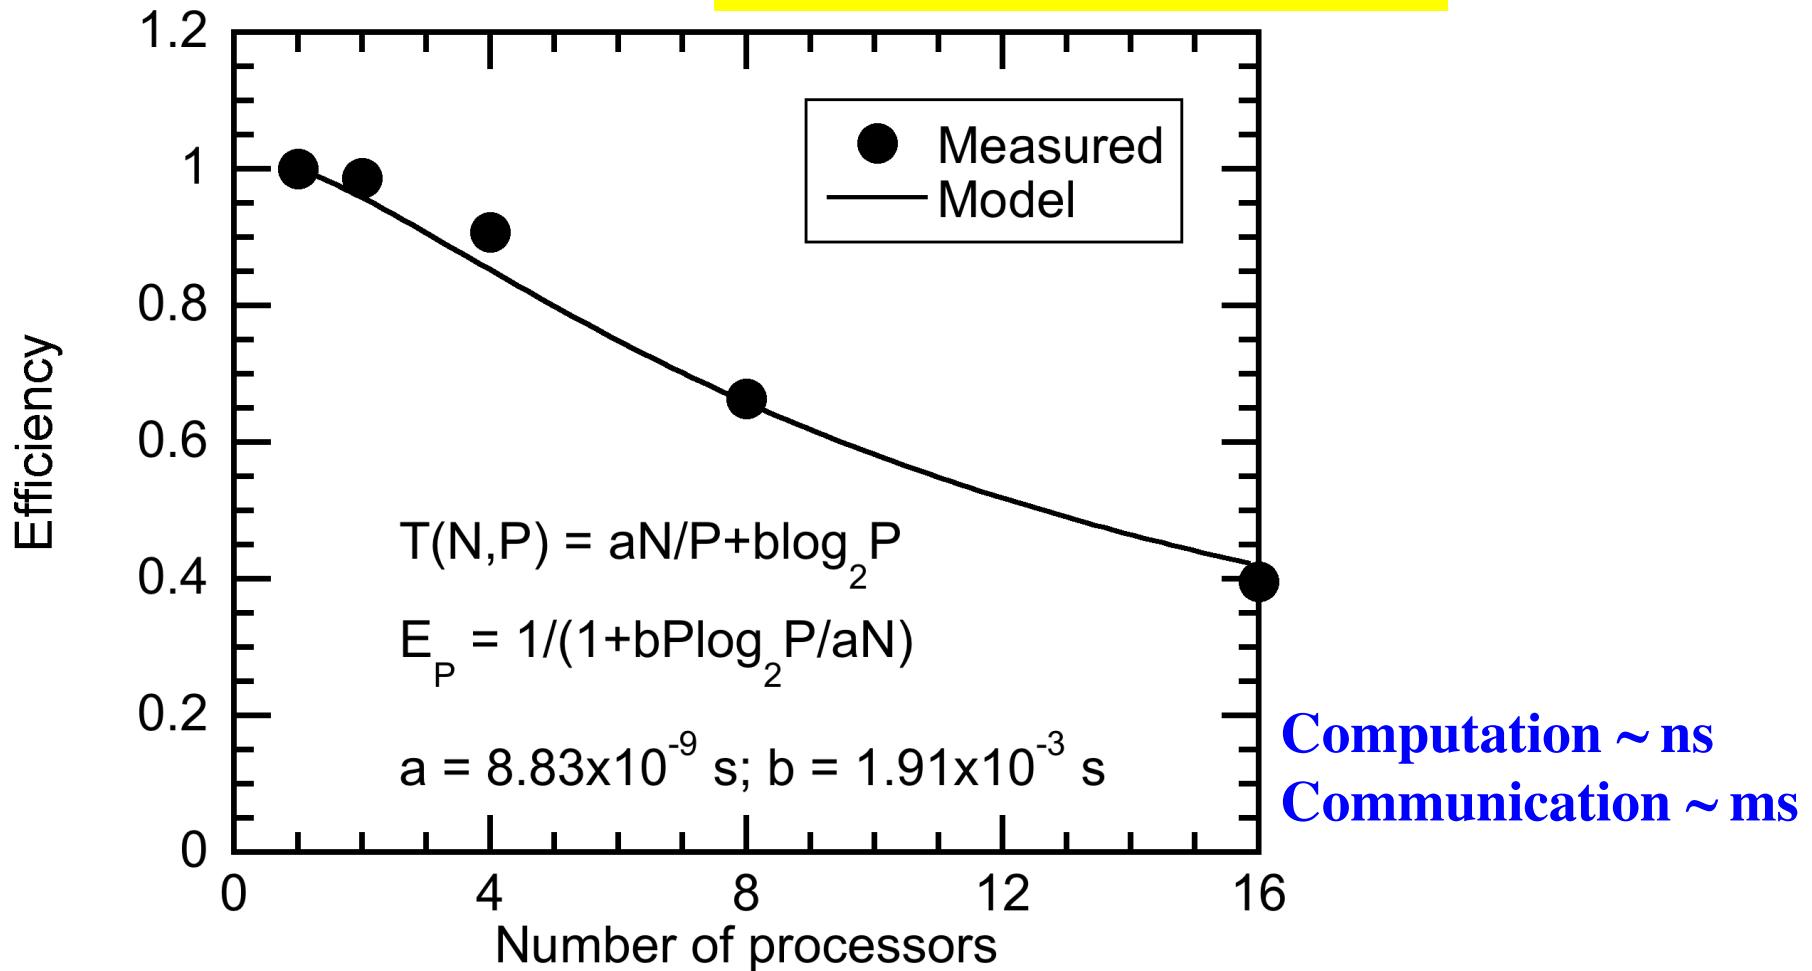


global_pi.c: $N = 10^7$, on HPC

Fixed Problem-Size Scaling

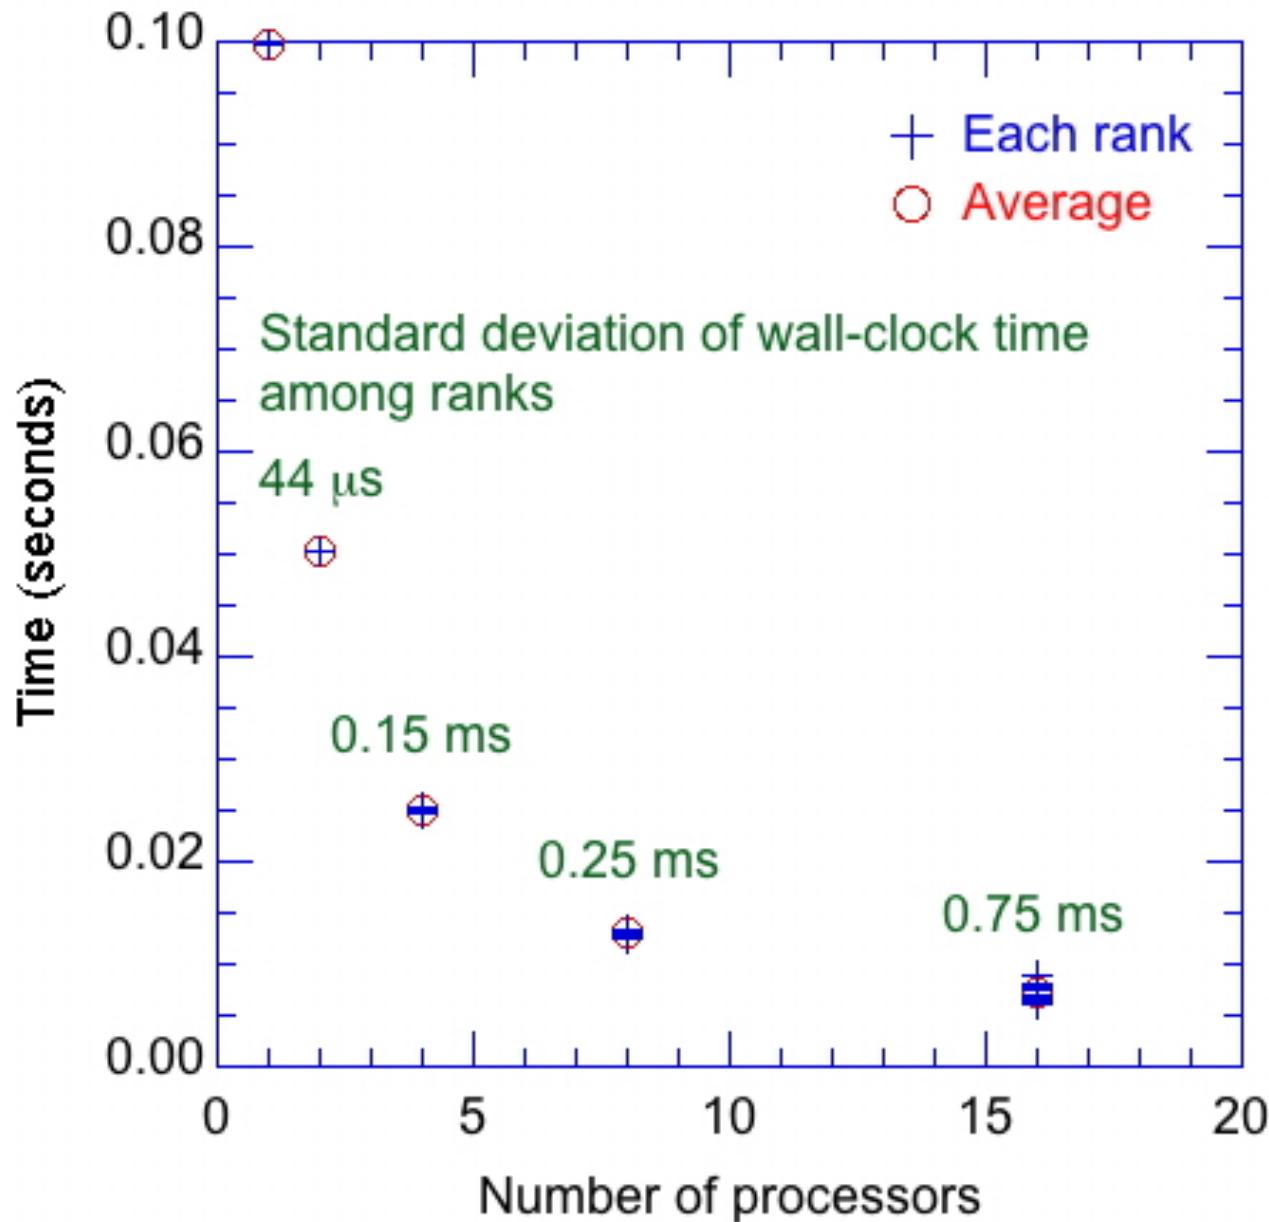
- Speedup model:

$$E_P = \frac{S_P}{P} = \frac{1}{1 + \frac{\beta}{\alpha} \frac{P \log P}{N}}$$



global_pi.c: $N = 10^7$, on HPC

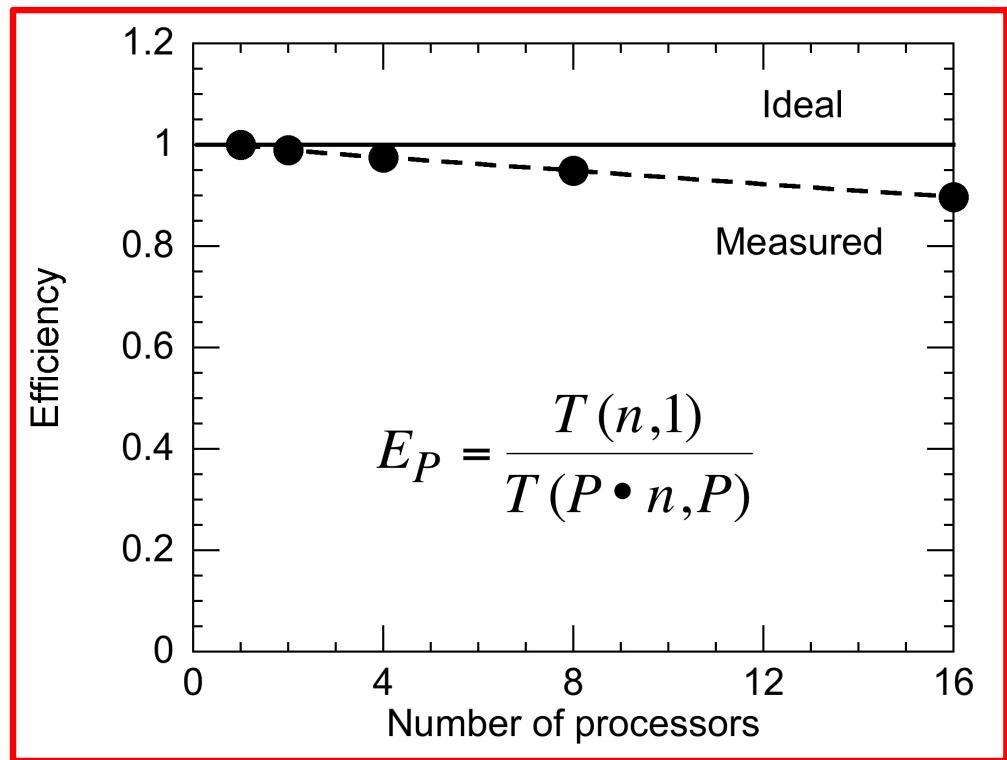
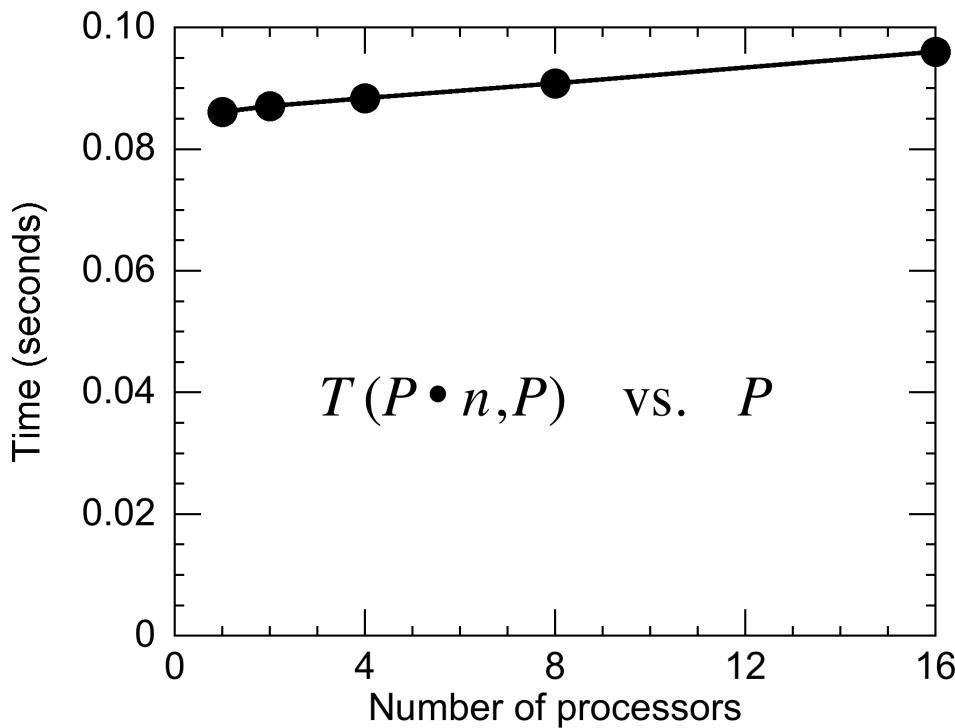
Runtime Variance among Ranks



Isogranular (Weak) Scaling

- $n = N/P = \text{constant}$
- Efficiency:

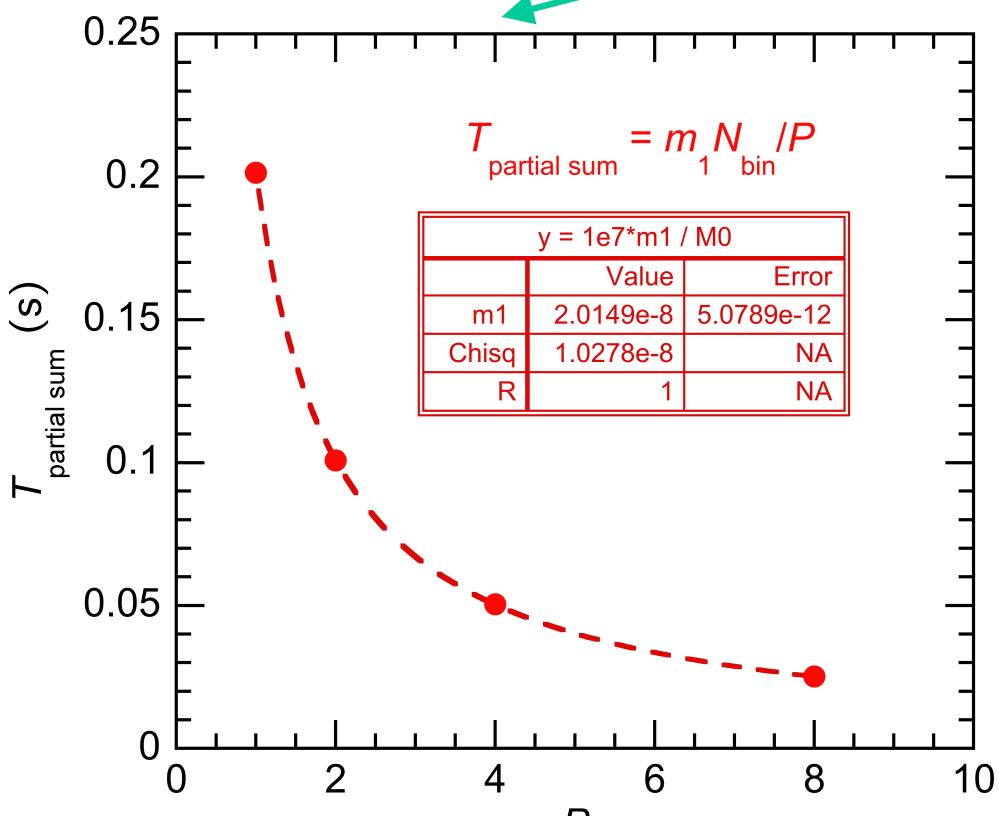
$$E_P = \frac{T(n, 1)}{T(nP, P)} = \frac{\alpha n}{\alpha n + \beta \log P} = \frac{1}{1 + \frac{\beta}{\alpha n} \log P}$$



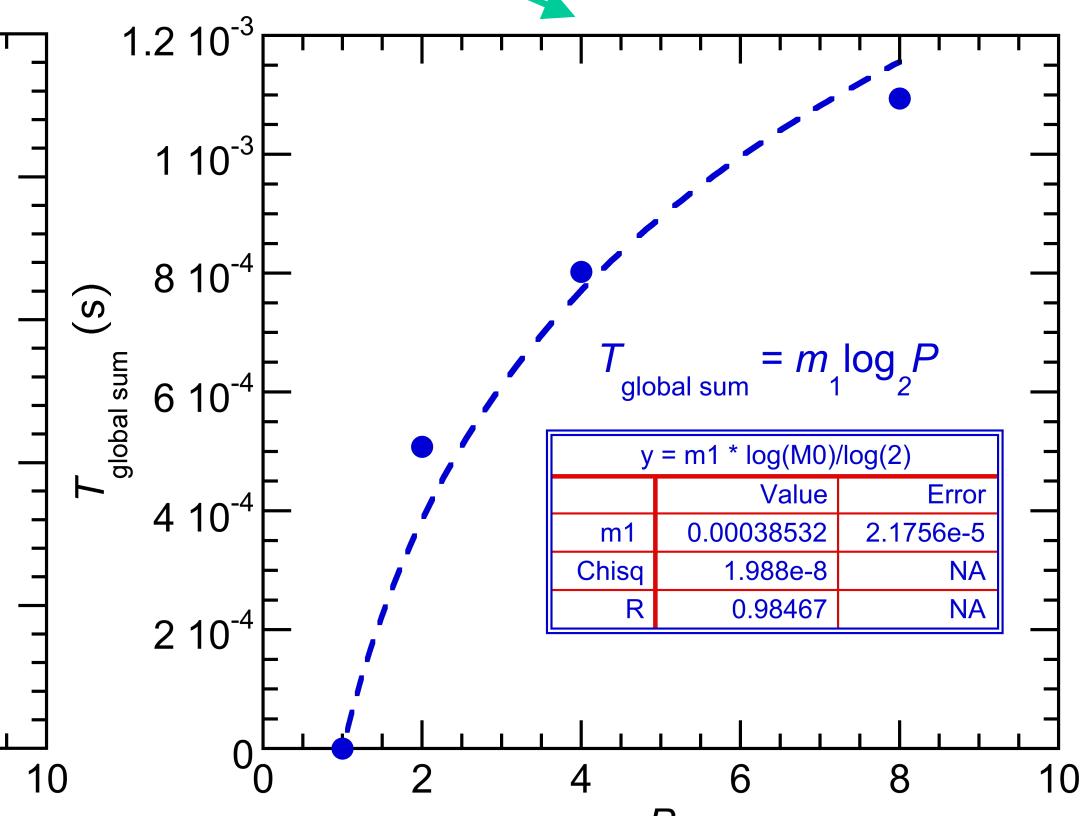
global_pi_iso.c: $N/P = 10^7$, on HPC

Fitting Fixed Problem-Size Scaling

$$T(N, P) = \alpha \frac{N}{P} + \beta \log_2 P$$



Computation ~ ns



Communication ~ ms

`global_pi_breakdown.c`: $N = 10^7$, on 3.0 GHz Xeon

Scalability Analysis: Example

Computer Physics Communications 219 (2017) 246–254

A derivation and scalable implementation of the synchronous parallel kinetic Monte Carlo method for simulating long-time dynamics

Hyekyung Suk Byun ^a, Mohamed Y. El-Naggar ^{a,b,c}, Rajiv K. Kalia ^{a,d,e,f}, Aiichiro Nakano ^{a,b,d,e,f,*},
Priya Vashishta ^{a,d,e,f}

$$T(N, P) = T_{\text{comp}}(N, P) + T_{\text{comm}}(N, P) + T_{\text{global}}(P) \\ = aN/P + b(N/P)^{2/3} + c \log P. \quad (6)$$

For isogranular scaling, the number of atoms per processor, $N/P = n$, is constant, and the isogranular parallel efficiency is

$$E_P = \frac{T(n, 1)}{T(nP, P)} = \frac{an}{an + bn^{2/3} + c \log P} \\ = \frac{1}{1 + \frac{b}{a}n^{-1/3} + \frac{c}{an} \log P}. \quad (7)$$

For fixed problem-size scaling, the global number of hemes, N , is fixed, and the speedup is given by

$$S_P = \frac{T(N, 1)}{T(N, P)} = \frac{aN}{aN/P + b(N/P)^{2/3} + c \log P} \\ = \frac{P}{1 + \frac{b}{a}\left(\frac{P}{N}\right)^{1/3} + \frac{c}{a}\frac{P \log P}{N}}, \quad (8)$$

and the parallel efficiency is

$$E_P = \frac{S_P}{P} = \frac{1}{1 + \frac{b}{a}\left(\frac{P}{N}\right)^{1/3} + \frac{c}{a}\frac{P \log P}{N}}. \quad (9)$$

