

Algorithms PART II: Partitioning and Divide & Conquer

HPC Fall 2012

Prof. Robert van Engelen






Overview

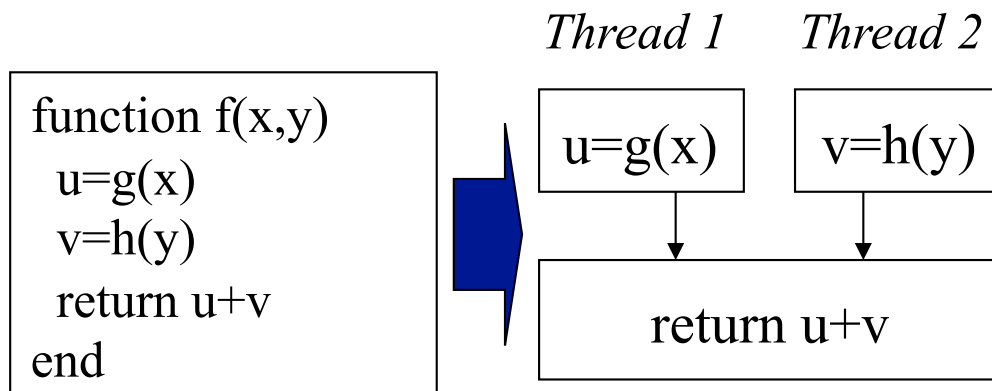
- Partitioning strategies
- Divide and conquer strategies
- Further reading

Partitioning Strategies



Block partitioning of a 2D domain

- *Data partitioning*
 - Perform *domain decomposition* to run parallel tasks on subdomains
 - “Scatter-compute-gather” where local computation may require communication and scatter/gather may involve computations



- *Task partitioning*
 - Decompose functions into independent subfunctions and execute the subfunctions in parallel



Partitioning Strategies

- *Partitioning strategy* (data partitioning):
 1. Break up a given problem into P subproblems
 2. Solve the P subproblems concurrently
 3. Collect and combine the P solutions

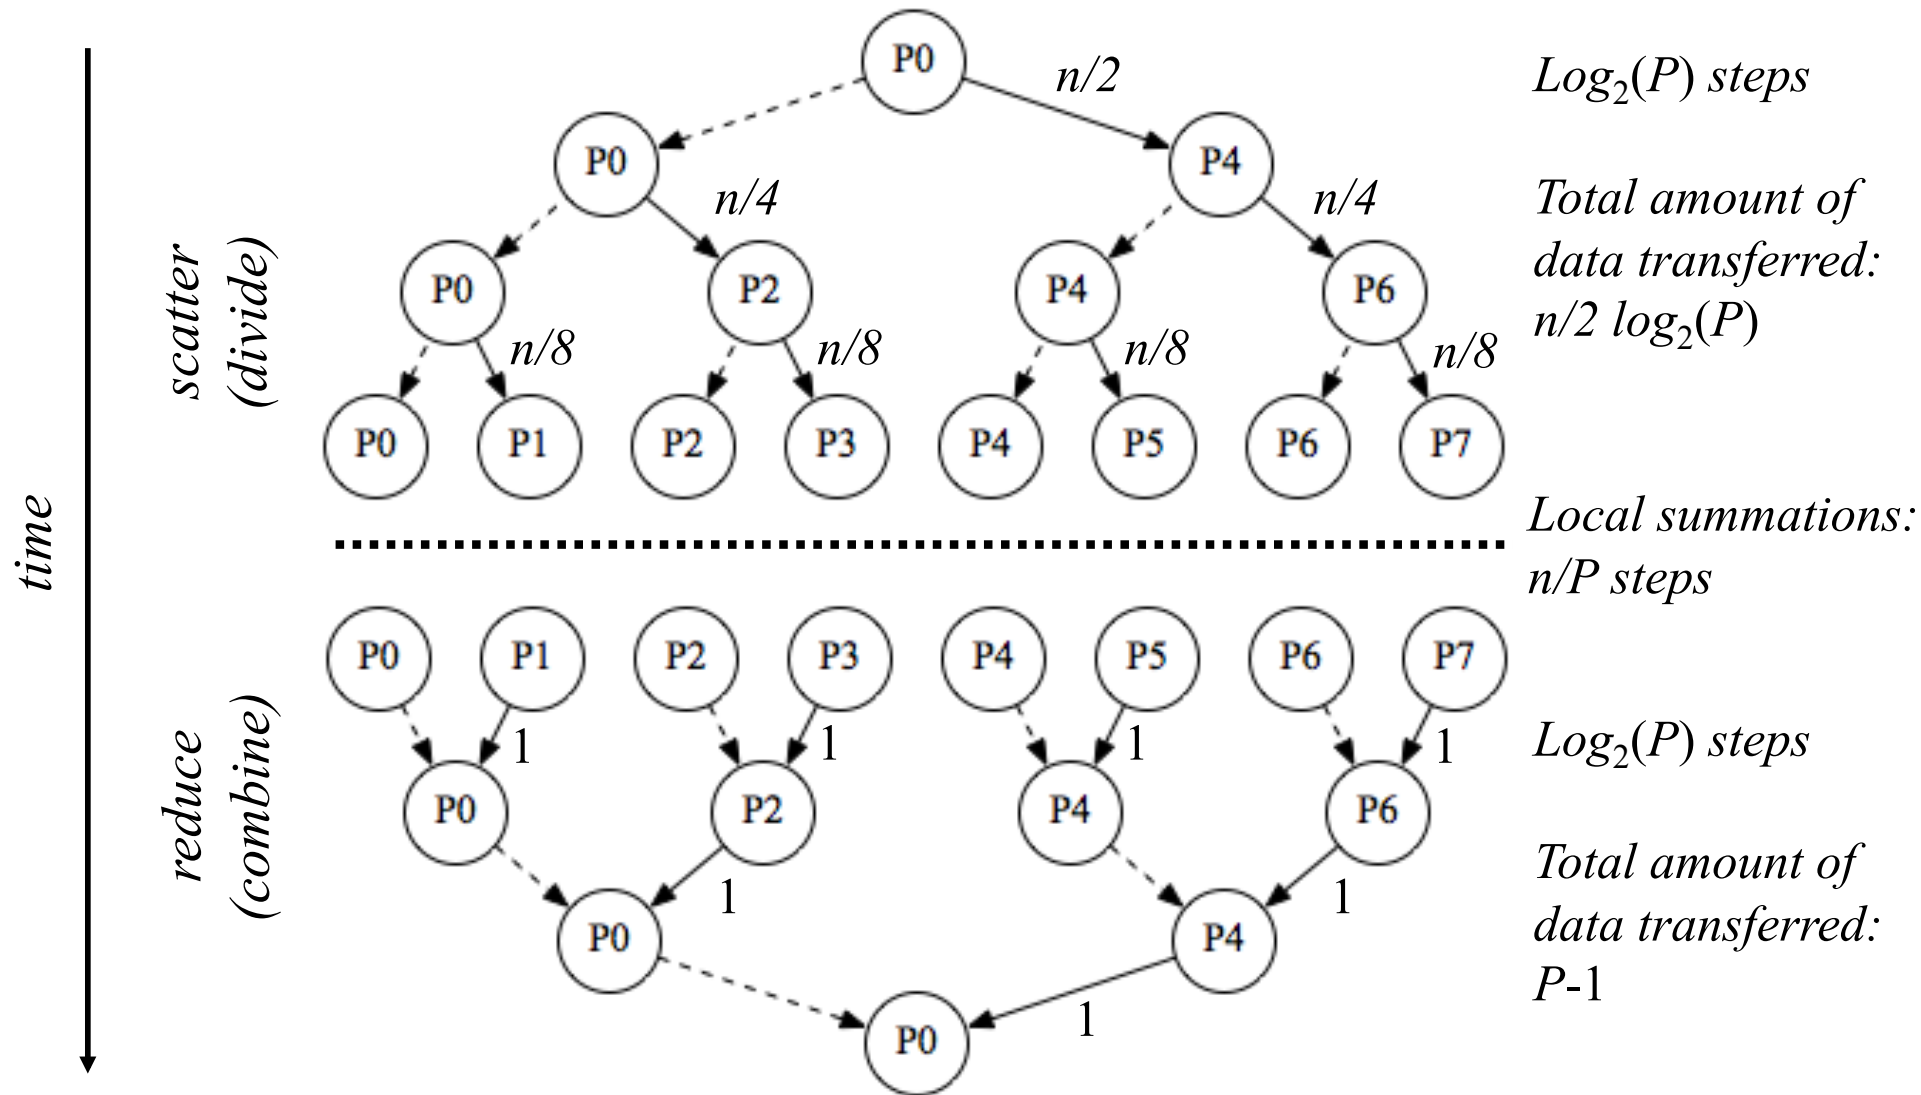
- Embarrassingly parallel
 - Is a simple form of data partitioning into independent subproblems without initial work and no communication between tasks (workers)



Partitioning Example 1: Summation

- Summation of n values $X = [x_1, \dots, x_n]$
 1. Divide X into P equally-sized sublists X_p , $p = 0, \dots, P-1$ and distribute the X_p sublists to the P processors
 2. The processors sum the local parts $s_p = \sum X_p$
 3. Combine the local sums $s = \sum s_p$
- Algorithms:
 1. Scatter list X using a scatter-tree
 2. Serial summation of parts
 3. Reduce local sums

Partitioning Example 1: Summation





Partitioning Example 1: Summation

■ Communication time

- Scatter:
$$t_{comm1} = \sum_{k=1..log_2(P)} (t_{startup} + 2^{-k}n t_{data})$$
$$= \log_2(P)t_{start} + n(P-1)/P t_{data}$$
- Reduce:
$$t_{comm2} = \log_2(P) (t_{start} + t_{data})$$
- Total:
$$t_{comm} = 2 \log_2(P)t_{start} + (n(P-1)/P + \log_2(P)) t_{data}$$


■ Computation time

- Local sum:
$$t_{comp1} = n/P$$
- Global sum:
$$t_{comp2} = \log_2(P)$$
- Total:
$$t_{comp} = n/P + \log_2(P)$$


■ Speedup, assuming $t_{startup} = 0$

- Sequential time: $t_s = n-1$
- Parallel time: $t_p = (n(P-1)/P + \log_2(P)) t_{data} + n/P + \log_2(P)$
- Speedup: $S_p = t_s/t_p = O(n / (n + \log(P)))$
- Best speedup w/o communication: $S_p = O(P/\log(P))$

General M-Ary Partitioning



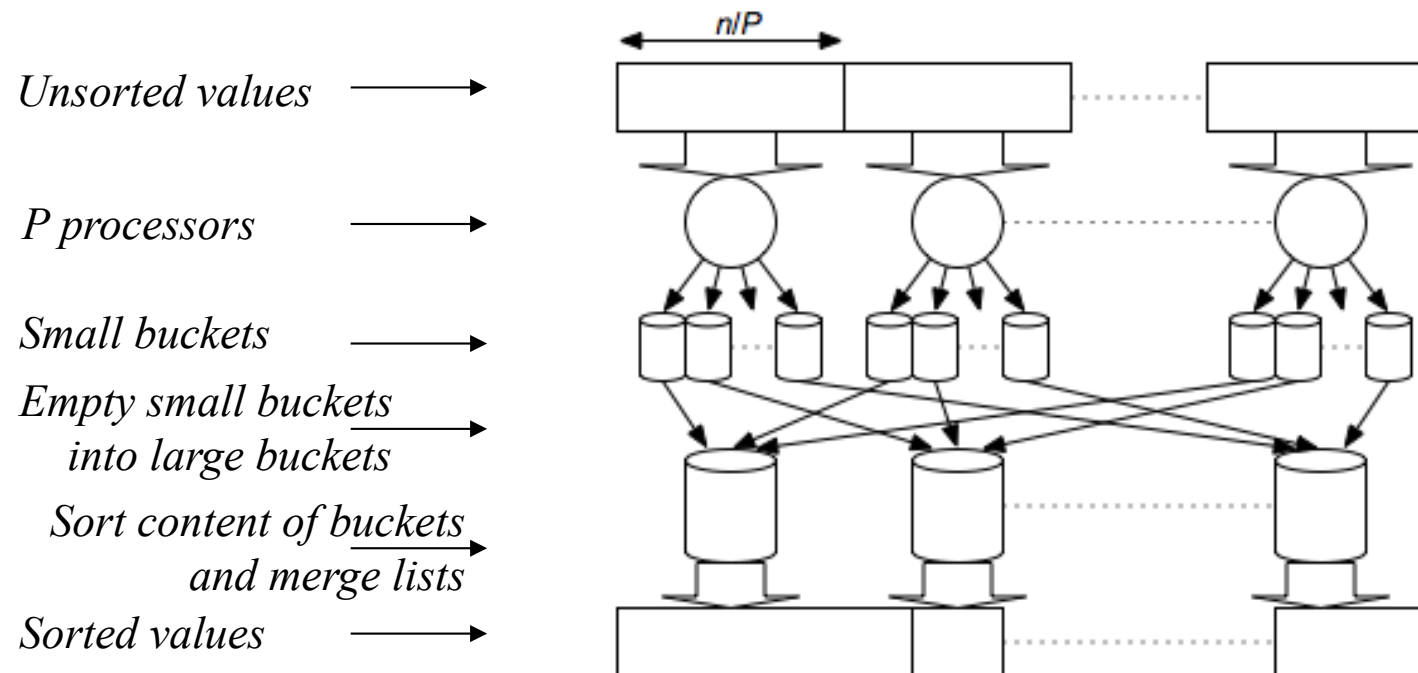
Example: partitioning an image, e.g. to compute histogram by parallel reductions (summations to count color pixels)



3-level 4-ary partitioning for $4^3 = 64$ processors

Partitioning Example 2: Parallel Bucket Sort

- Bucket sort of values $[x_1, \dots, x_n]$ bounded within a range $x_i \in [l_0 \dots h_i]$
 1. Partition the n values in n/P segments
 - 2a. Sort each segment into P small buckets (local computation)
 - 2b. Send content of small buckets to P large buckets
 3. Sort P large buckets and merge lists





Partitioning Example 2: Parallel Bucket Sort

Input: list X of length n with minimum value L and maximum U

Output: sorted list X

```
def function bucket( $x$ ) =  $P * (x - L) / (U - L)$  ;
```

```
scatter list  $X$  to local  $X_p$  lists each of size  $n/P$ 
```

```
forall processors  $p = 0, \dots, P-1$ 
```

```
  for  $i = 0, \dots, n/P-1$ 
```

```
     $x = X_p[i]$ 
```

```
    put  $x$  into small bucket  $b_p[bucket(x)]$ 
```

```
all-to-all of small buckets  $b_p$  into large buckets  $B_p$ 
```

```
sort values in  $B_p[0, \dots, P-1]$  using a sequential sort algorithm
```

```
gather  $X$  from  $B_p$  into a merged sorted list
```



Partitioning Example 2: Parallel Bucket Sort

- Communication time (assuming uniform distribution in X)
 - Scatter: $t_{comm1} = \log_2(P)t_{startup} + n(P-1)/P t_{data}$
 - All-to-all: $t_{comm2} = (P-1)(t_{startup} + n/P^2 t_{data})$
 - Gather: $t_{comm3} = \log_2(P)t_{startup} + n(P-1)/P t_{data}$
- Computation time (assuming uniform distribution in X)
 - Small bucket sort: $t_{comp1} = n/P$
 - Large bucket sort: $t_{comp2} = n/P \log_2(n/P)$
- Speedup
 - Sequential time: $t_s = n \log_2(n/P)$ (with P buckets)
 - Parallel time: $t_p = 2 \log_2(P)t_{startup} + 2 n(P-1)/P t_{data} + (P-1)(t_{startup} + n/P^2 t_{data}) + n/P (1 + \log_2(n/P))$
 - Speedup w/o communication: $S_p = O(P)$



Partitioning Example 3: Barnes Hut Algorithm

$$F = \frac{Gm_1m_2}{r^2}$$

Direction of the force between
two bodies at points p and q

$$\vec{F} = \frac{Gm_p m_q}{r^2} \left(\frac{\vec{p} - \vec{q}}{r} \right)$$


$$F = ma$$

$$\vec{F}^t = \frac{m(\vec{v}^{t+\frac{1}{2}} - \vec{v}^{t-\frac{1}{2}})}{\Delta t}$$

$$\vec{v}^{t+\frac{1}{2}} = \vec{v}^{t-\frac{1}{2}} + \frac{\vec{F}^t \Delta t}{m}$$

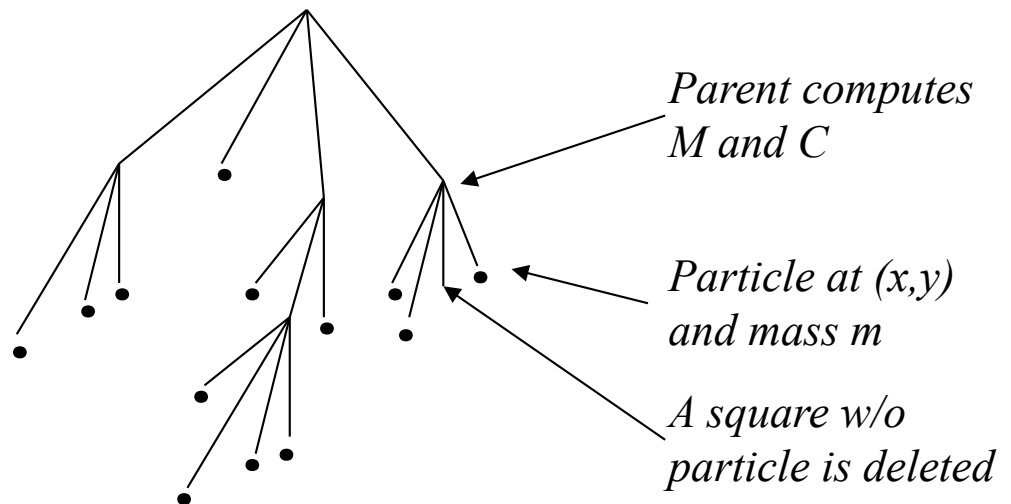
$$\vec{x}^{t+1} = \vec{x}^t + \vec{v}^{t+\frac{1}{2}} \Delta t$$

Partitioning Example 3: Barnes Hut Algorithm



Particles in 2D space

Quadtree



Mass of parent is sum of masses of children

$$M = \sum_{i=0}^3 m_i$$

Center of mass

$$C = \frac{1}{M} \sum_{i=0}^3 m_i c_i$$



Partitioning Example 3: Barnes Hut Algorithm

```
for (t = 0; t < tmax; t++)  
{  
    Build_tree();  
    Compute_Total_Mass_Center();  
    Compute_Force();  
    Update_Positions();  
}
```

Sequential time is $O(n \log n)$

Assuming $P = n$ then $t_p = O(\log P)$

$$C = \frac{1}{M} \sum_{i=0}^3 m_i c_i \quad (*)$$

$$\vec{F} = \frac{G m_p m_q}{r^2} \left(\frac{\vec{p} - \vec{q}}{r} \right) \quad (**)$$

```
Compute_Force()  
{  
    for (i = 0; i < n; i++)  
        Compute_Tree_Force(i, root)  
}  
Compute_Tree_Force(i, node)  
{  
    if (box at node contains one particle)  
        F = force using eq (**)  
    else  
    {  
        r = distance from i to C (*) of box  
        D = size of box at node  
        if (D/r < theta)  
            F = force using eq (**) with total M  
        else  
            for (all children c of box)  
                F = F + Compute_Tree_Force(i, c);  
    }  
    return F;  
}
```



Divide and Conquer

- Divide and conquer strategy (definition by JáJá 1992)
 1. Break up a given problem into independent subproblems
 2. Solve the subproblems ***recursively*** and concurrently
 3. Collect and combine the solutions into the overall solution
- In contrast to the partitioning strategy, divide and conquer uses *recursive partitioning* with concurrent execution to divide the problem down into independent subproblems
- In deeper levels of recursion the number of active processors may increase or decrease



Divide & Conquer Example 1: Parallel Recursive Matmul

- Block matrix multiplication in recursion by decomposing matrix in 2×2 submatrices and computing the submatrices recursively

```
Mat matmul(Mat A, Mat B, int s)
{ if (s == 1)
  C = A * B;
  else
  { s = s/2;
    P0 = matmul(Ap,p, Bp,p, s);
    P1 = matmul(Ap,q, Bq,p, s);
    P2 = matmul(Ap,p, Bp,q, s);
    P3 = matmul(Ap,q, Bq,q, s);
    P4 = matmul(Aq,p, Bp,p, s);
    P5 = matmul(Aq,q, Bq,p, s);
    P6 = matmul(Aq,p, Bp,q, s);
    P7 = matmul(Aq,q, Bq,q, s);
    Cp,p = P0 + P1;
    Cp,q = P2 + P3;
    Cq,p = P4 + P5;
    Cq,q = P6 + P7;
  }
  return C;
}
```

P0...P7 computed in parallel


Can be computed in parallel

- Level of parallelism increases with deepening recursion
- Suitable for shared memory systems



Divide and Conquer Example 2: Parallel Convex Hull Algorithm


- The *planar convex hull* of a set of points $S = \{p_1, p_2, \dots, p_n\}$ of $p_i = (x, y)$ coordinates is the smallest convex polygon that encompasses all points S on the x - y plane






Divide and Conquer Example 2: Parallel Convex Hull Algorithm

- The *upper convex hull* spans points $\{q_1, \dots, q_s\} \subseteq S$ from point q_1 with minimum x to q_s with maximum x
- The *convex hull* = *upper convex hull* + *lower convex hull*
- Problem:
 - Given points $S = \{p_1, \dots, p_n\}$ such that $x(p_1) < x(p_2) < \dots < x(p_n)$, construct the upper convex hull in parallel



Divide and Conquer Example 2: Parallel Convex Hull Algorithm

- Parallel convex hull:
 1. Divide the x -sorted points S into sets S_1 and S_2 of equal size
 2. Compute upper convex hull recursively on S_1 and S_2
 3. Combine $UCH(S_1)$ and $UCH(S_2)$ by computing the upper common tangent a to b to form $UCH(S)$





Divide and Conquer Example 2: Parallel Convex Hull Algorithm

- Base case of recursion: two points, which are returned as $UCH(S)$
- The line segment (a,b) can be computed sequentially in $O(\log n)$ time with $n = |UCH(S_1) + UCH(S_2)|$ using a binary search method
- Line segments can be implemented as linked list of points, thus $UCH(S_1)$ and $UCH(S_2)$ can be connected using one pointer change of a to point to b in $O(1)$ time
- Parallel convex hull time complexity recurrence relation:
$$T(n) \leq T(n/2) + a \log n$$

with solution:
$$T(n) = O(\log^2 n)$$
- Parallel convex hull operations recurrence relation:
$$W(n) \leq 2W(n/2) + b n$$

with solution:
$$W(n) = O(n \log n)$$

which is cost optimal, since sequential algorithm is $O(n \log n)$



Divide and Conquer Example 3: First-Order Linear Recurrences

- First-order linear recurrence

$$y_1 = b_1$$

$$y_i = a_i y_{i-1} + b_i \quad 2 \leq i \leq n$$

- Example applications:

- Prefix sum $y_i = \sum_{j=1..i} b_j$ is a special case of a first-order linear recurrence with $a_i = 1$ (the multiplicative unit element)

- n-th order polynomial evaluation using Horner's rule

$$p(x) = (((b_1 x + b_2) x + b_3) x + \dots + b_{n-1}) x + b_n$$

is a special case of a first-order linear recurrence with $a_i = x$

- Solving a bi-diagonal system $\mathbf{B}\mathbf{y} = \mathbf{c}$,

let

$$a_i = -l_i/d_i$$

$$b_i = c_i/d_i$$

then solve linear recurrence

to obtain solution \mathbf{y}

$$\begin{pmatrix} d_1 & & & & \\ l_2 & d_2 & & & \\ & l_3 & d_3 & & \\ & & \dots & \dots & \\ & & & l_n & d_n \end{pmatrix} \begin{pmatrix} y_1 \\ y_2 \\ y_3 \\ \dots \\ y_n \end{pmatrix} = \begin{pmatrix} c_1 \\ c_2 \\ c_3 \\ \dots \\ c_n \end{pmatrix}$$



Divide and Conquer Example 3: First-Order Linear Recurrences

- Rewrite $y_i = a_i y_{i-1} + b_i$ into $y_i = a_i (a_{i-1} y_{i-2} + b_{i-1}) + b_i$
- This equation defines a linear recurrence of size $n/2$ for even index i

$$\begin{aligned} z_1 &= b_1' \\ z_i &= a_i' z_{i-1} + b_i' \quad 2 \leq i \leq n/2 \end{aligned}$$

1. Let


$$\begin{aligned} a_i' &= a_{2i} a_{2i-1} \\ b_i' &= a_{2i} b_{2i-1} + b_{2i} \end{aligned}$$

2. Solve z_i recursively

3. For $1 \leq i \leq n$ set

$$\begin{aligned} y_i &= z_{i/2} && \text{if } i \text{ is even} \\ y_i &= a_i z_{(i-1)/2} + b_i && \text{if } i \text{ is odd } > 1 \\ y_i &= b_1 && \text{if } i = 1 \end{aligned}$$

Divide and Conquer Example 3: First-Order Linear Recurrences



$\log_2 n$ recursive steps

■ Parallel algorithm:

```
linrecsolve(a[], b[], y[], n)
{
    if (n==1)
    { y[1] = b[1];
      return;
    }
    forall (i = 1 to n/2)
    { a_new[i] = a[2*i]*a[2*i-1];
      b_new[i] = a[2*i]*b[2*i-1]+b[2*i];
    }
    linrecsolve(a_new, b_new, z, n/2);
    forall (i = 1 to n)
    { if (i == 1)
      { y[1] = b[1];
        else if (even(i))
        { y[i] = z[i/2];
          else
          { y[i] = a[i]*z[(i-1)/2]+b[i];
        }
      }
    }
}
```

$$\begin{aligned}
 &b_1 \\
 &b_1' = a_2 b_1 + b_2 \\
 &b_1'' = a_2' b_1' + b_2' \\
 &\quad = ((a_2 b_1 + b_2) a_3 + b_3) a_4 + b_4 \\
 &b_1''' = a_2'' b_1'' + b_2'' \\
 &\quad = ((a_2' b_1' + b_2') a_3' + b_3') a_4' + b_4' \\
 &\quad = (((a_2 b_1 + b_2) a_3 + b_3) a_4 + b_4) a_5 + b_5) a_6 + b_6) a_7 + b_7) a_8 + b_8
 \end{aligned}$$



Divide and Conquer Example 4: Triangular Matrix Inversion

- Consider $\mathbf{Ax} = \mathbf{b}$ with $n \times n$ triangular matrix \mathbf{A}

$$\begin{pmatrix} a_{11} & & & & \\ a_{21} & a_{22} & & & \\ a_{31} & a_{32} & a_{33} & & \\ \dots & \dots & \dots & \dots & \\ a_{n1} & a_{n2} & \dots & \dots & a_{nn} \end{pmatrix}$$

- Partition \mathbf{A} into $(n/2) \times (n/2)$ blocks

$$\begin{bmatrix} \mathbf{A}_1 & \\ \mathbf{A}_2 & \mathbf{A}_3 \end{bmatrix}$$

- Then \mathbf{A}^{-1} is given by

$$\begin{bmatrix} \mathbf{A}_1^{-1} & \mathbf{0} \\ -\mathbf{A}_3^{-1}\mathbf{A}_2\mathbf{A}_1^{-1} & \mathbf{A}_3^{-1} \end{bmatrix}$$



Divide and Conquer Example 4: Triangular Matrix Inversion

- Parallel algorithm:
 1. Divide A into A_1 , A_2 , A_3
 2. Recursively compute inverses of A_1 and A_3 in parallel
 3. Multiply $-A_3^{-1}A_2A_1^{-1}$ and combine with A_1^{-1} and A_3^{-1} to get A^{-1}
- Time complexity is given by the recurrence relation
$$T(n) = T(n/2) + c n$$
with $P=n^2$ processors to compute $-A_3^{-1}A_2A_1^{-1}$ in $O(n)$ operations in parallel, thus $T(n) = O(n)$ time

Divide and Conquer Example 5: Banded Triangular Systems

- Consider $\mathbf{Ax} = \mathbf{b}$ with banded matrix \mathbf{A} with $m=3$

$$\begin{pmatrix}
 a_{11} & & & & & & & & \\
 a_{21} & a_{22} & & & & & & & \\
 a_{31} & a_{32} & a_{33} & & & & & & \\
 & a_{42} & a_{43} & a_{44} & & & & & \\
 & & a_{53} & a_{54} & a_{55} & & & & \\
 & & & a_{64} & a_{65} & a_{66} & & & \\
 & & & & a_{75} & a_{76} & a_{77} & & \\
 & & & & & a_{86} & a_{87} & a_{88} & \\
 & & & & & & a_{97} & a_{98} & a_{99}
 \end{pmatrix}
 \rightarrow
 \begin{pmatrix}
 \begin{array}{|c|c|c|}
 \hline
 a_{11} & & \\
 a_{21} & a_{22} & \\
 a_{31} & a_{32} & a_{33} \\
 \hline
 \end{array}
 & & \\
 & \begin{array}{|c|c|c|}
 \hline
 a_{42} & a_{43} & a_{44} \\
 & a_{53} & a_{54} & a_{55} \\
 & a_{64} & a_{65} & a_{66} \\
 \hline
 \end{array}
 & \\
 & & \begin{array}{|c|c|c|}
 \hline
 & a_{75} & a_{76} & a_{77} \\
 & & a_{86} & a_{87} & a_{88} \\
 & & & a_{97} & a_{98} & a_{99} \\
 \hline
 \end{array}
 \end{pmatrix}$$

- Define block diagonal \mathbf{D} and inverse \mathbf{D}^{-1}

$$\mathbf{D} = \begin{pmatrix}
 \mathbf{A}_{11} & & & \\
 & \mathbf{A}_{22} & & \\
 & & \dots & \\
 & & & \dots & \mathbf{A}_{n/m, n/m}
 \end{pmatrix}
 \quad
 \mathbf{D}^{-1} = \begin{pmatrix}
 \mathbf{A}_{11}^{-1} & & & \\
 & \mathbf{A}_{22}^{-1} & & \\
 & & \dots & \\
 & & & \dots & \mathbf{A}_{n/m, n/m}^{-1}
 \end{pmatrix}$$



Divide and Conquer Example 5: Banded Triangular Systems

- Compute $\mathbf{d} = \mathbf{D}^{-1}\mathbf{b}$ and $\mathbf{B} = \mathbf{D}^{-1}\mathbf{A}$ where $\mathbf{B}_{i,i-1} = \mathbf{A}_{ii}^{-1}\mathbf{A}_{i,i-1}$

$$\mathbf{d} = \mathbf{D}^{-1}\mathbf{b} = \begin{pmatrix} \mathbf{d}_1 \\ \mathbf{d}_2 \\ \dots \\ \dots \\ \mathbf{d}_{n/m} \end{pmatrix} \quad \mathbf{B} = \mathbf{D}^{-1}\mathbf{A} = \begin{pmatrix} \mathbf{I}_m & & & & \\ \mathbf{B}_{21} & \mathbf{I}_m & & & \\ & \mathbf{B}_{32} & \mathbf{I}_m & & \\ & & \dots & \dots & \\ & & & \mathbf{B}_{n/m,n/m-1} & \mathbf{I}_m \end{pmatrix}$$

- Solve first-order linear recurrence on $m \times m$ matrices $\mathbf{B}_{i,i-1}$

$$\mathbf{x}_1 = \mathbf{d}_1$$

$$\mathbf{x}_i = -\mathbf{B}_{i,i-1} \mathbf{x}_{i-1} + \mathbf{d}_i \quad 2 \leq i \leq n/m$$

- Parallel time $O(m + m \log(n/m))$ with $P=nm$ processors
 - Compute all \mathbf{A}_{ii}^{-1} (each requiring $O(m)$ operations) in parallel with parallel matrix inversion algorithm
 - Compute all $\mathbf{B}_{i,i-1} = \mathbf{A}_{ii}^{-1}\mathbf{A}_{i,i-1}$ in $O(m)$ operations in parallel
 - Recurrence depth is $\log_2(n/m)$, each step has $O(m)$ operations



Divide and Conquer Example 6: LU of Tridiagonal Matrix

- Consider tridiagonal matrix LU decomposition

$$\begin{pmatrix} a_1 & c_1 & & & \\ b_2 & a_2 & c_2 & & \\ & b_3 & a_3 & c_3 & \\ & & \dots & \dots & \dots \\ & & & b_n & a_n \end{pmatrix} = \begin{pmatrix} 1 & & & & \\ l_2 & 1 & & & \\ & l_3 & 1 & & \\ & & \dots & \dots & \\ & & & l_n & 1 \end{pmatrix} \begin{pmatrix} d_1 & u_1 & & & \\ & d_2 & u_2 & & \\ & & d_3 & u_3 & \\ & & & \dots & \dots \\ & & & & d_n \end{pmatrix}$$

- The LU decomposition $\mathbf{A} = \mathbf{L} \mathbf{U}$ satisfies

$$a_1 = d_1$$

$$c_i = u_i$$

$$a_i = d_i + l_i u_{i-1}$$

$$b_i = l_i d_{i-1}$$

thus

$$d_1 = a_1$$

$$d_i = a_i - l_i u_{i-1} = a_i - u_{i-1} b_i / d_{i-1} = [a_i d_{i-1} - b_i c_{i-1}] / d_{i-1}$$



Divide and Conquer Example 6: LU of Tridiagonal Matrix

- Let

$$\mathbf{R}_1 = \begin{bmatrix} a_1 & 0 \\ 1 & 0 \end{bmatrix} \quad \mathbf{R}_i = \begin{bmatrix} a_i & -b_i c_{i-1} \\ 1 & 0 \end{bmatrix} \quad \mathbf{T}_i = \mathbf{R}_i \mathbf{R}_{i-1} \dots \mathbf{R}_1$$

- From the Möbius transformation we have

$$d_i = \frac{\begin{bmatrix} 1 \\ 0 \end{bmatrix}^T \mathbf{T}_i \begin{bmatrix} 1 \\ 1 \end{bmatrix}}{\begin{bmatrix} 0 \\ 1 \end{bmatrix}^T \mathbf{T}_i \begin{bmatrix} 1 \\ 1 \end{bmatrix}}$$

- Algorithm:

- Set up matrices \mathbf{R}
- Solve first-order linear recurrence (prefix sum) of \mathbf{T}
- Compute d_i
- From the solution of d_i compute $l_i = b_i/d_{i-1}$



Further Reading

- [PP2] pages 106-131
- [PSC] pages 321-337